

CMPE 200  
Computer Architecture & Design

## **Lecture 4.** **Memory Hierarchy (4)**

Haonan Wang




SAN JOSÉ STATE  
UNIVERSITY

# Cache Policies: Cases

- **Allocation policy: do we allocate a block in cache for the missed data?**
- **Read policies:**
  - Read Hit: this is what we want. Only one data read from the cache.
  - Read Miss: needs to fetch from lower level, but just write to the register once after that
    - read-allocate (with replacement policy) vs. no-read-allocate (i.e., cache bypassing)
- **Write policies (only for the data cache): consistency & performance tradeoffs**
  - Write Hit: behavior and number of writes depends on write policy
    - Write-through vs. write-back vs. write-evict
  - Write Miss: needs to first read from lower level, then apply write policies
    - Write-allocate (with replacement policy): Write-through vs. Write-back
    - No-write-allocate (bypassing): Write-evict

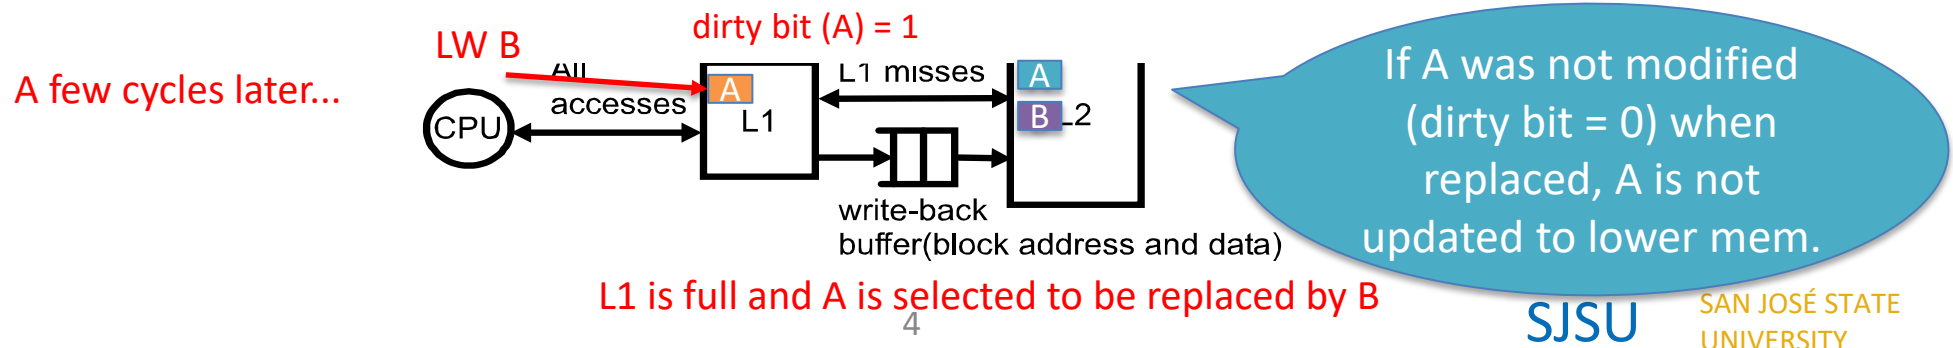
# Write Policy: Write-back

- **Write-back: inconsistent with lower level**
  - The value is written only to the cache line.
  - The modified (dirty) cache line is written to the lower level only when it is evicted.
    - 1 dirty bit is needed for each cache line.
  - A write-back buffer to update lower level with evicted dirty blocks
    - Must write a full block at this point since we do not know which word is modified
  - Example: assume cache block containing word address A is initially in the cache



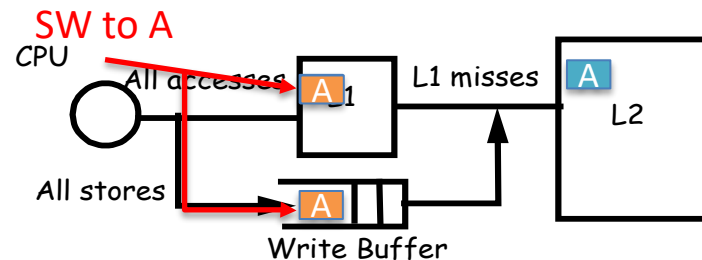
# Write Policy: Write-back

- **Write-back: inconsistent with lower level**
  - The value is written only to the cache line.
  - The modified (dirty) cache line is written to the lower level only when it is evicted.
    - 1 dirty bit is needed for each cache line.
  - A write-back buffer to update lower level with evicted dirty blocks
    - Must write a full block at this point since we do not know which word is modified
  - Example: assume cache block containing word address A is initially in the cache




# Write Policy: Write-through

- **Write-through: consistent with lower level**
  - The value is written to both the cache line and the lower-level memory.
    - No need to write a full block to the lower level (e.g., only write 4B for a SW)
  - Write buffer can be used so cache does not need to wait for the write to lower level
    - stall only if the write buffer is full
  - Example: assume cache block containing word address A is initially in the cache



# Write Policy: Write-through

- **Write-through: consistent with lower level**
  - The value is written to both the cache line and the lower-level memory.
    - No need to write a full block to the lower level (e.g., only write 4B for a SW)
  - Write buffer can be used so cache does not need to wait for the write to lower level
    - stall only if the write buffer is full
  - Example: assume cache block containing word address A is initially in the cache



# Write Policy: Write-evict

- **Write-evict (no-write): no consistency issue**
  - The value is written only to the lower level (i.e., bypassing).
  - No need to fetch the data first if it is missing in the cache
  - If the block containing the write address exists in the cache, evict it.
  - The write only needs to send the data specified by the store instruction (e.g., 4 Bytes).
  - Write buffer can be used so cache does not need to wait for the write to lower level
  - More expensive when there is read after write

# Cache Policies: Cases - Revisit

- **Allocation policy: do we allocate a block in cache for the missed data?**
- **Read policies:**
  - Read Hit: this is what we want. Only one data read from the cache.
  - Read Miss: needs to fetch from lower level, but just write to the register once after that
    - read-allocate (with replacement policy) vs. no-read-allocate (i.e., cache bypassing)
      - + write policy of evicted data
- **Write policies (only for the data cache): consistency & performance tradeoffs**
  - Write Hit: behavior and number of writes depends on write policy
    - Write-through vs. write-back vs. write-evict
  - Write Miss: needs to first read from lower level, then apply write policies
    - Write-allocate (with replacement policy): Write-through vs. Write-back
      - + write policy of evicted data
    - No-write-allocate (bypassing): Write-evict



# Cache Miss Behavior Analysis

- **Read miss:**
  - **+Write-through:** evict victim block + fetch block from lower level
  - **+Write-back:** evict victim block + write back evicted block if dirty + fetch block from lower level
  - **+No-write:** find victim block + fetch block from lower level
- **Write miss:**
  - Write-allocate:
    - **+Write-through:** evict victim block + fetch block from lower level + store word to block + store word to lower level
    - **+Write-back:** evict victim block + write back evicted block if dirty + fetch block from lower level + store word to block
  - No-write-allocate: **+Write-evict:** store word to lower level directly

# Other Cache Policy details

- **Some caches support multiple policies**
  - Can achieve different delay/bandwidth tradeoffs or even dynamically change
- **Write hit and write miss can use different policies**
  - E.g., write-hit write-evict + write-miss write-allocate-write-back
- **Different cache levels can also use different policies**
  - E.g., L1 use write-evict + L2 use write-back
- **Cache can also be completely bypassed (not the same as write-evict)**
  - E.g., bypassing L1

# Reduce Miss Rate (1): Code Optimization

- **Misses occur if sequentially accessed array elements come from different cache blocks**
- **Code optimizations → No hardware change**
  - Rely on programmers or compilers
- **Examples:**
  - Loop interchange: In nested loops, outer loop becomes inner loop and vice versa
  - Loop blocking: partition large array into smaller blocks, thus fitting the accessed array elements into cache size

# Loop Interchange

**Assume: cache line size = 5 words**

What addresses are accessed in each iteration of inner loop?


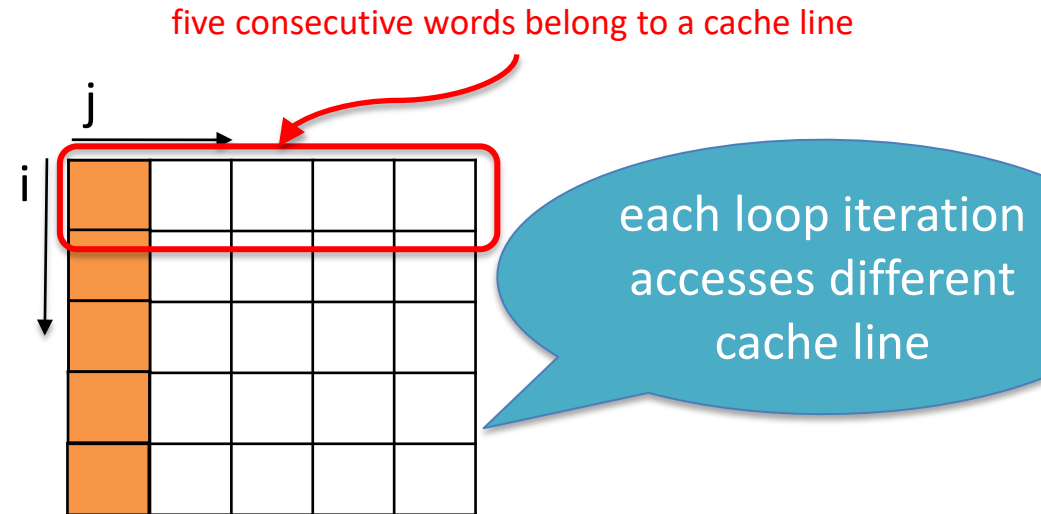
```
/* Before */  
for (j=0; j<5; j++)  
  for (i=0; i<5; i++)  
    x[i][j] = 2*x[i][j]
```

**Column-major ordering**



**Row-major ordering**


```
/* After */  
for (i=0; i<5; i++)  
  for (j=0; j<5; j++)  
    x[i][j] = 2*x[i][j]
```



# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.



# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```

The second matrix may always encounter misses because individual data are from different cache blocks




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

# Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```

If N is large, the cache line that has been accessed is likely to be replaced before next access




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

# Loop Blocking

In the block-based matrix multiplication:

```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```


Matrix is computed in  
a smaller block unit  
→ higher locality



# Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



# Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



# Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



# Loop Blocking

In the block-based matrix multiplication:

```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```






# Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



# Loop Blocking

In the block-based matrix multiplication:

```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



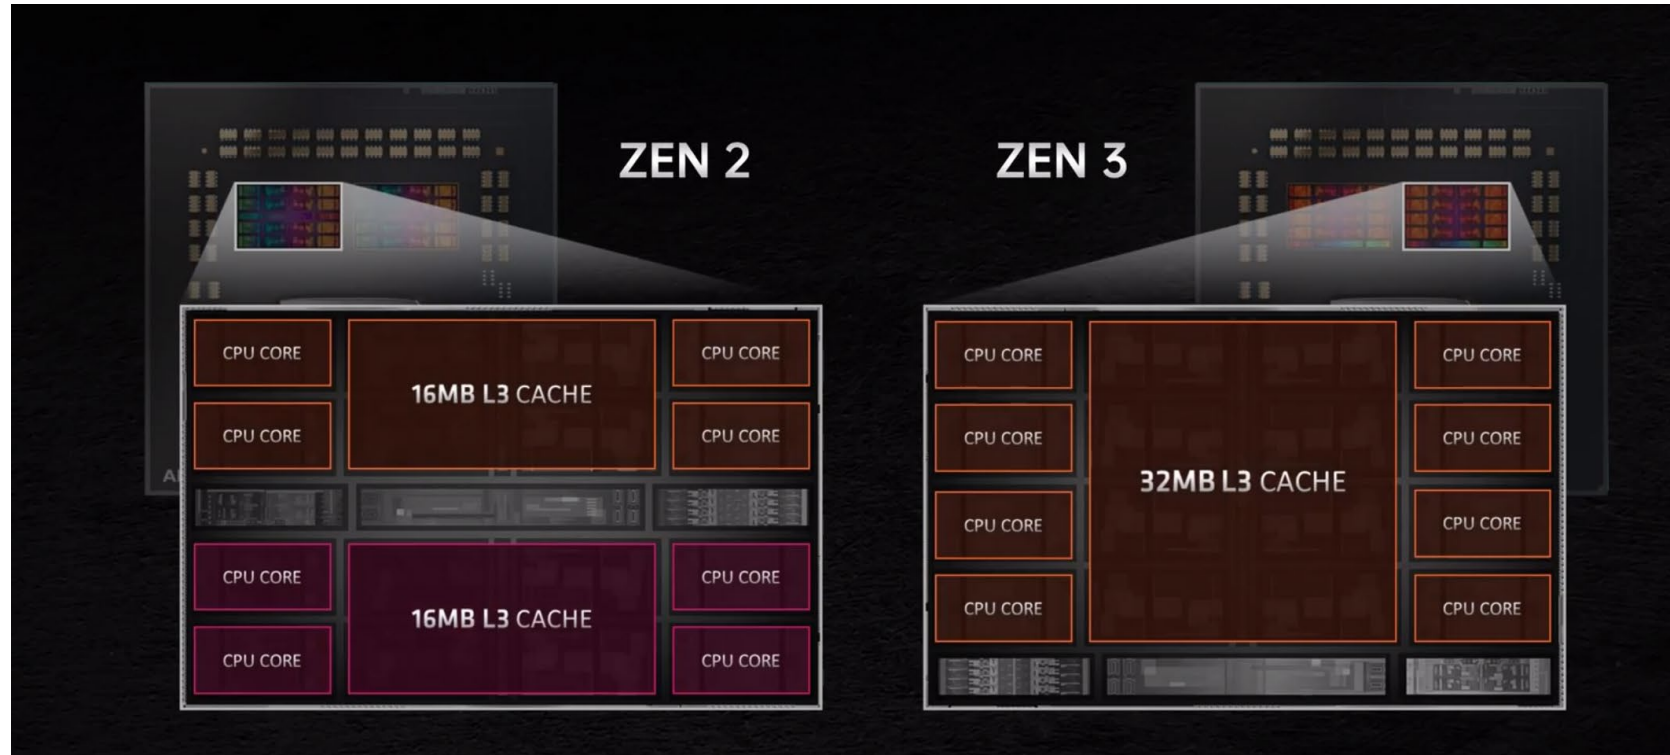
# Cache Miss Classification: The 3 C's

- **Compulsory (cold) Misses**
  - On the 1<sup>st</sup> reference to a block
  - Related to # blocks accessed by a code, not related to the configuration of a cache
- **Capacity Misses**
  - The program's working set size exceeds the cache capacity
- **Conflict Misses**
  - Multiple memory blocks map to the same set in set-associative caches


# Reduce Miss Rate (2): Reduce the 3 C's

- **Increase Cache Size**
  - Reduce miss for: capacity miss, conflict miss
  - But has many limitations
- **Increase Associativity (cache size unchanged)**
  - Reduce miss for: conflict miss
  - But may increase access latency
- **Increase Cache Block Size (cache size unchanged)**
  - Reduce miss for: compulsory
  - But may increase miss penalty (more data will be evicted and fetched)
  - Very large blocks could increase miss rate (i.e., increase capacity miss)

# Increase Cache Size




# Miss Rate vs Block Size vs Cache Size




- Cache size – the larger the better
- Block size – tradeoffs **Why?**

# Block Size Tradeoff



$$\text{Average Memory Access Time} = \text{Hit Time} + \text{Miss Penalty} \times \text{Miss Rate}$$

# Benefits of Set Associative Caches



Largest gains when going from direct mapped to 2-way (20%+ reduction in miss rate)



# Reduce Miss Rate (3): Policies

- **Use appropriate policies**
  - LRU + write-back works for most applications
  - Many GPUs use write-hit write-evict + write-miss no-write-allocate
  - Can even dynamically change policy according to profiling
- **Use more advanced policies**
  - E.g., ML based

# Reduce Miss Rate (4): Multi-level Caches

- Having a unified L2 cache (i.e., it holds both instructions and data) and in some cases even a unified L3 cache
- L1 cache should focus on **minimizing hit time** in support of a shorter clock cycle
- Secondary cache(s) should focus on **reducing miss rate** to reduce the penalty of long main memory access times

# Split vs. Unified Caches

- **Split L1I\$ and L1D\$: to optimize  $t_{\text{hit}}$** 
  - Low capacity/associativity/block size (reduce  $t_{\text{hit}}$ )
  - Minimize structural hazards
  - Can optimize L1I\$ for wide output and no writes
- **Unified L2: to optimize hit rate**
  - $t_{\text{hit}}$  is less important due to less frequent accesses & long delay already
  - High capacity/associativity/block size (reduce  $\%_{\text{miss}}$ )
    - Unused instr capacity can be used for data (fewer capacity misses)
    - Instr / data conflicts (small to no increase for conflict misses in a large cache)
  - Instr / data structural hazards are rare (would take a simultaneous L1I\$ and L1D\$ miss)

# Reduce Miss Rate (5): Prefetching

- **Hardware prefetching**
  - Fetch blocks into the cache proactively (speculatively)
  - Key is to anticipate the upcoming miss addresses accurately
  - Relies on having unused memory bandwidth available
- **A simple case is to use **next block** prefetching**
  - Miss on address  $X \rightarrow$  anticipate next reference miss on  $X + \text{block-size}$
  - Works well for instructions (sequential execution) and for arrays of data
- **Need to initiate prefetches sufficiently in advance**
- **If prefetched instruction/data is not used, the cache is polluted with unnecessary data (possibly evicting useful data)**

# Other than Miss Rates

- **Cache bypassing (e.g., L1D)**
  - Reduce latency if L1 misses
- **Sector cache**
  - Fetch only a portion of the cache line
  - Not the same as reducing cache line size
- **Value prediction**
  - Using history value to predict future values
  - Roll back if predict wrong
  - Or accept a certain quality loss if the application is error tolerant

# Let us Conclude

What are the operations in cache for write-miss when using write-back?

evict victim block + write back evicted block if dirty + fetch block from lower level + store word to block

List two ways of code optimizations to improve locality?

Loop interchange, Loop blocking

# Conclusion Time

---

What are some ways to reduce cache misses?

Software methods, reduce 3C's, changing policies, multi-level caches, prefetching, value prediction ...

SAN JOSÉ STATE UNIVERSITY *powering* SILICON VALLEY

