

CMPE 200
Computer Architecture & Design

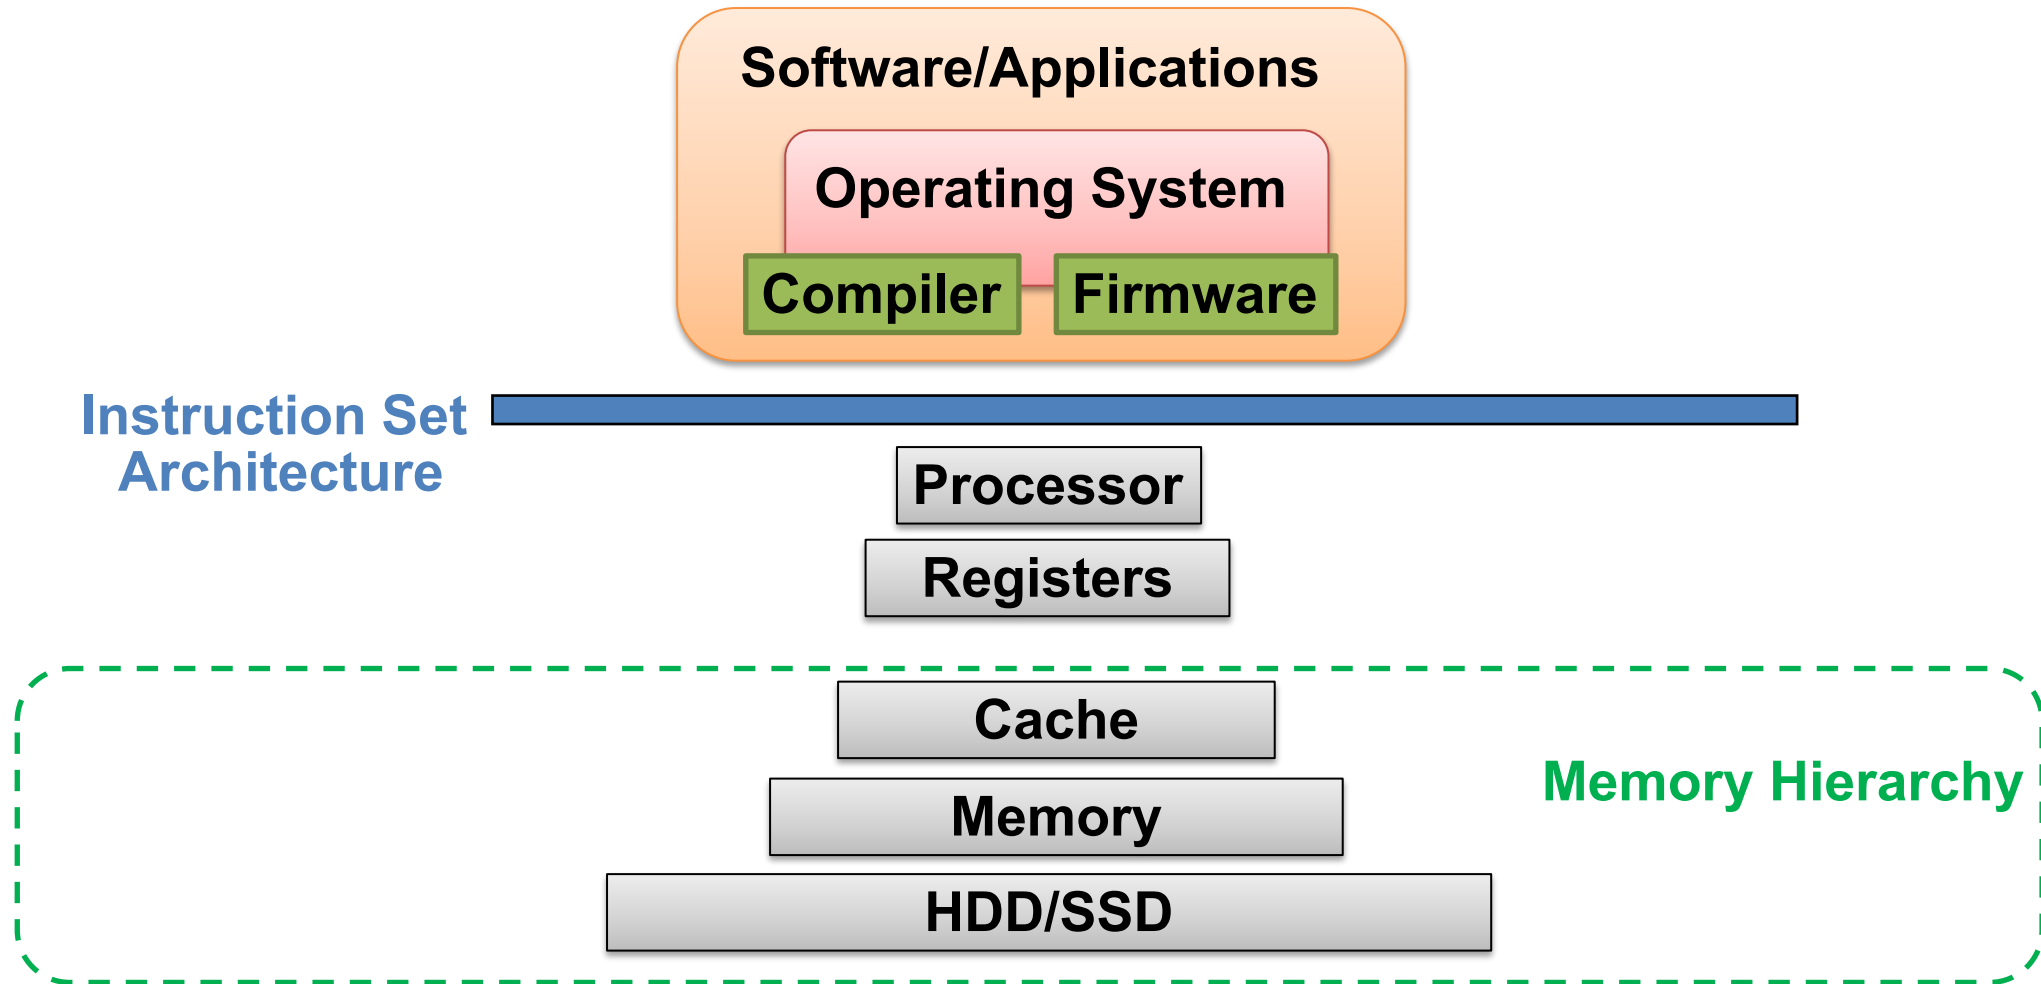
Lecture 4. **Memory Hierarchy**

Haonan Wang



SAN JOSÉ STATE
UNIVERSITY

Computer Architecture Overview



Memory Hierarchy

Example: a C program that reads two integer values from “file.txt” file and prints the sum of them.

```
#include <stdio.h>
#include <string.h>

int numbers[2];

void myfunction(void)
{
    FILE *fp;
    int size = 2;
    int sum = 0;

    /* Open file for reading */
    fp = fopen("mynumbers.txt", "r");

    /* Read and display data */
    fread(numbers, sizeof(int), size, fp);
    fclose(fp);
    sum = numbers[0] + numbers[1];
    printf("Sum = %d\n", sum);
}

int main (void) {
    myfunction();
    return(0);
}
```

Processor (CPU)

Understands
and executes
each line of the
code.

Uses fast on-
chip memories

Memory (DRAM)

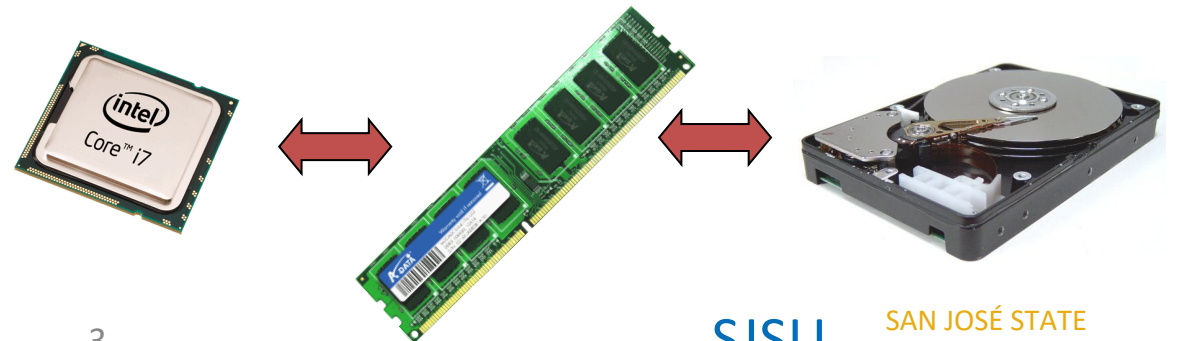
Provides
operands to
CPU

(*fp, size, sum,
numbers[2])

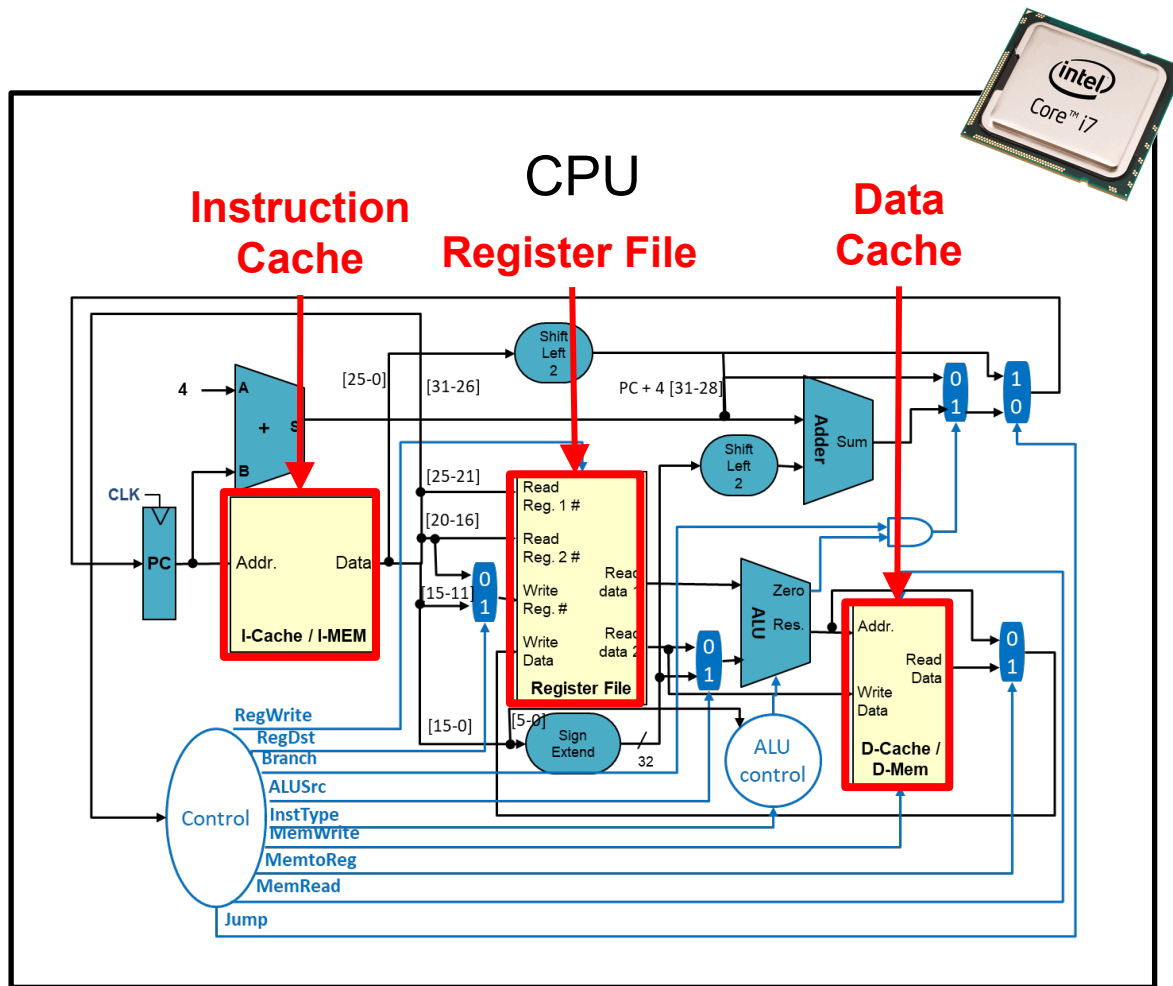
Storage (HDD)

Provides file
inputs and
program code


(file.txt)



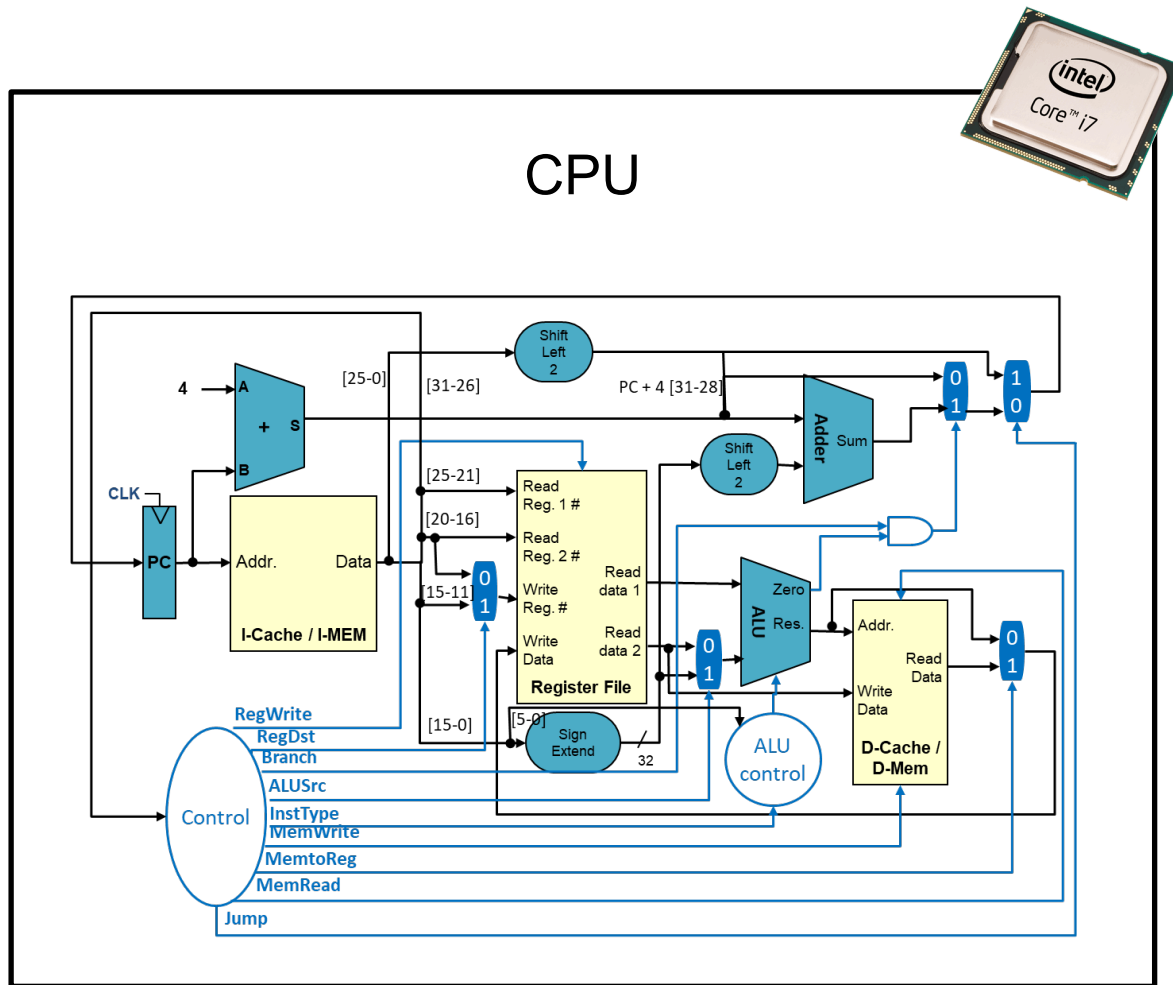
Memory Hierarchy



- **On-chip Memories (Memories inside of CPU)**
 - Register file, Caches
 - Small but Fast




Memory Hierarchy



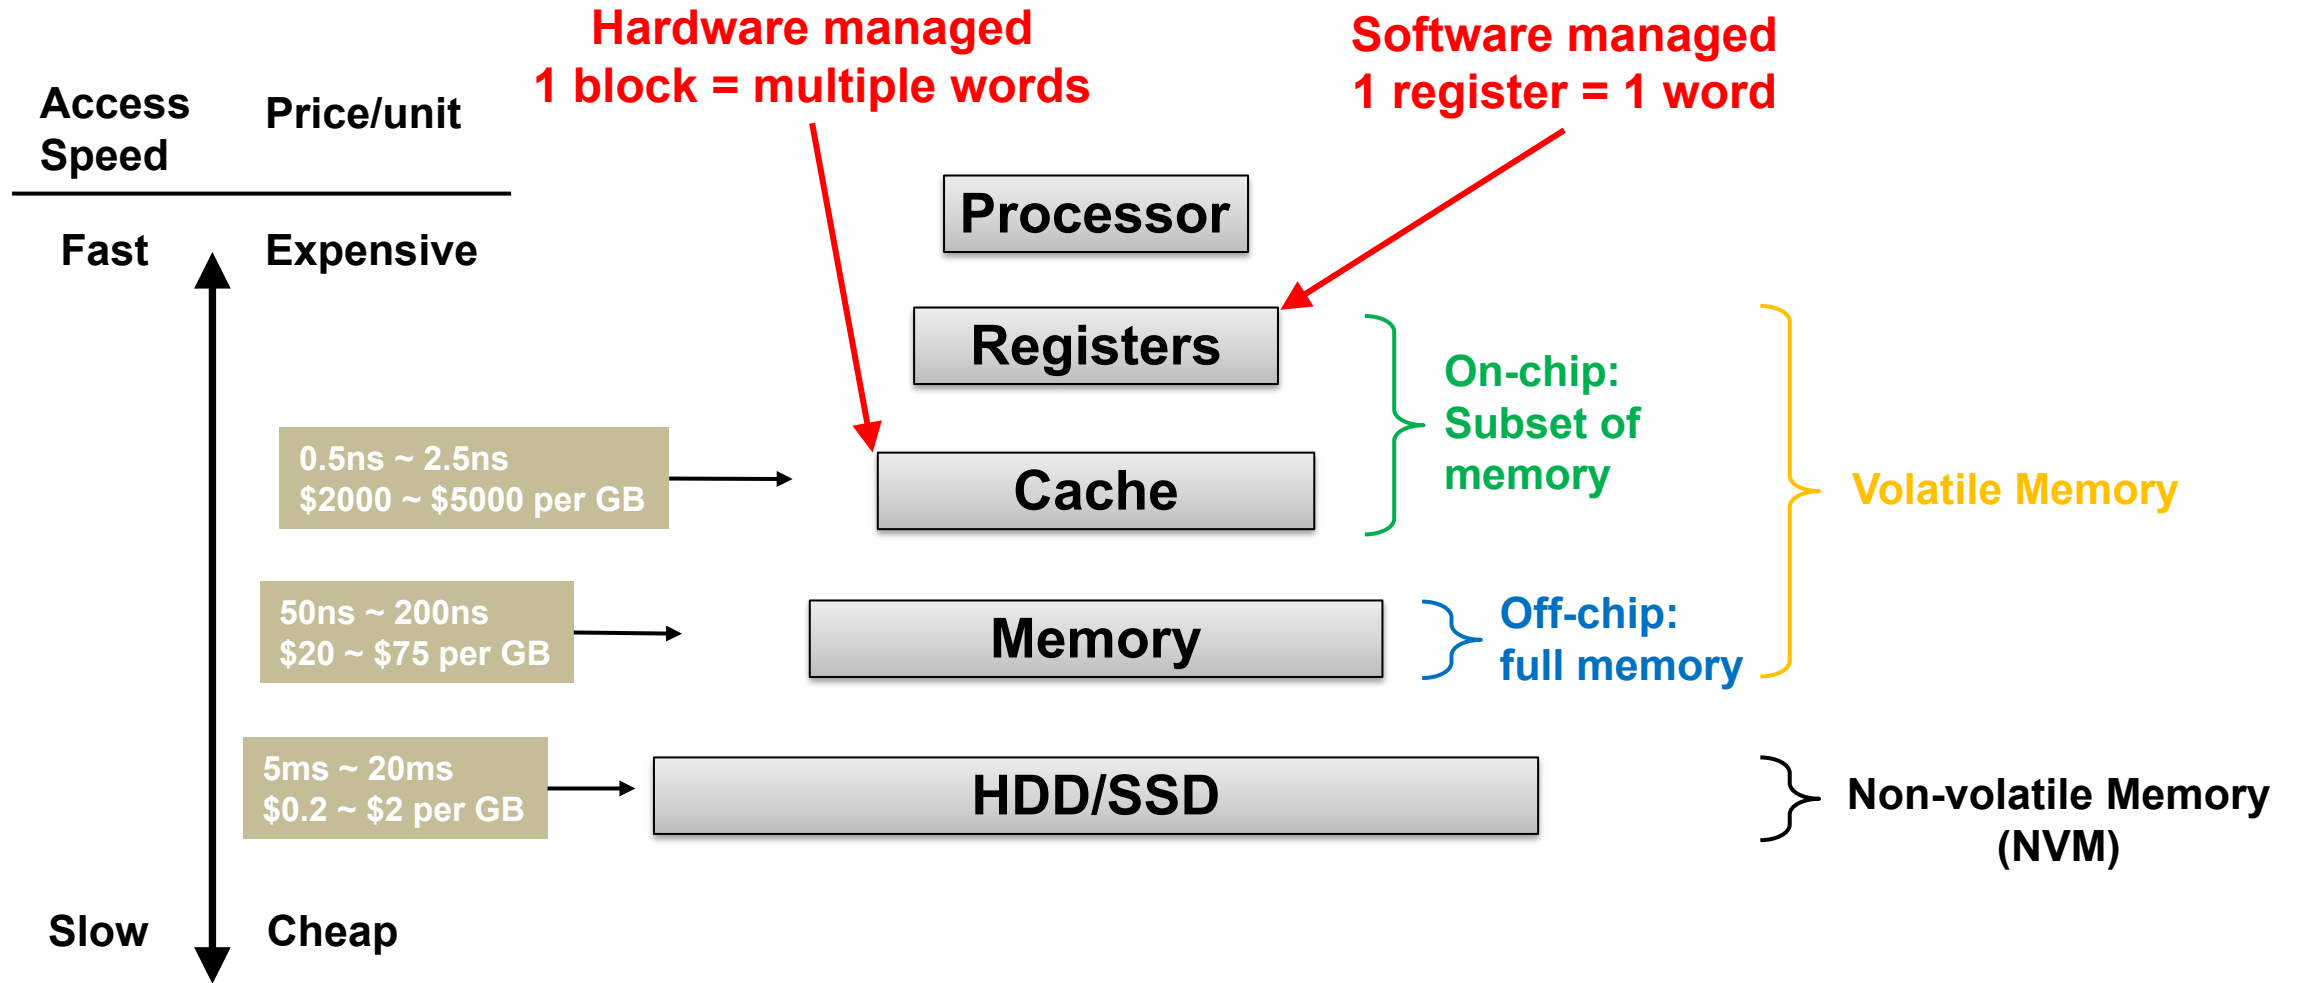
- **Off-chip Memories (Memories outside CPU)**
 - System Memory, Storage
 - Large but Slow

System Memory

Storage



Memory Hierarchy



Memories in Your PC


- **Windows**

- This PC → Properties
- cmd window → wmic
- 3rd party tool like CPU-Z

Windows edition

Windows 10 Education

© 2017 Microsoft Corporation. All rights reserved.



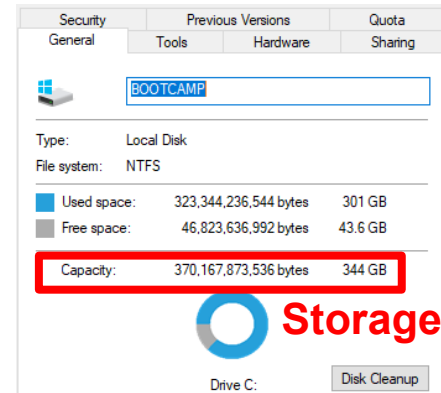
System

Processor: Intel(R) Core(TM) i7-4870HQ CPU @ 2.50GHz 2.50 GHz

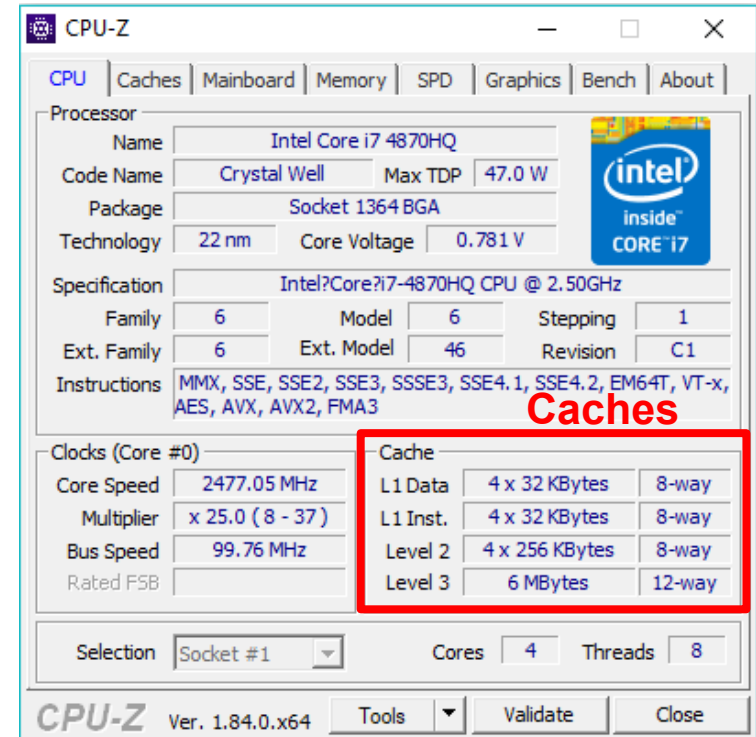
Installed memory (RAM): 16.0 GB

System type: 64-bit Operating System, x64-based processor

System Memory



Storage



Caches

Cache		
L1 Data	4 x 32 KBytes	8-way
L1 Inst.	4 x 32 KBytes	8-way
Level 2	4 x 256 KBytes	8-way
Level 3	6 MBytes	12-way


- **Linux**

- lscpu
- cat /proc/cpuinfo
- etc.


Why?

- **Two Types of Locality:**

- **Temporal Locality** (Locality in Time): If an address is referenced, it tends to be referenced again (e.g., loops, variable reuse)
- **Spatial Locality** (Locality in Space): If an address is referenced, neighboring addresses tend to be referenced (e.g., array, stack, etc.)



The “Memory Wall”




The Memory Hierarchy Goal

- **How do we create a memory system that gives the illusion of being large, cheap and fast (most of the time)?**
 - With hierarchy
 - try the fast parts first -- most of the time, this works well
 - if not, move the data so it works well the next time
 - With parallelism
 - use multiple identical parts operating simultaneously
 - for large quantities of data, this works well
- **Example – keep a subset of the data in fast memory**
- **Example – 1-byte-wide memory \rightarrow 4 \times 1-byte-wide memory \rightarrow 4-byte-wide memory**
 - load word takes 4 memory cycles vs. 1 memory cycle

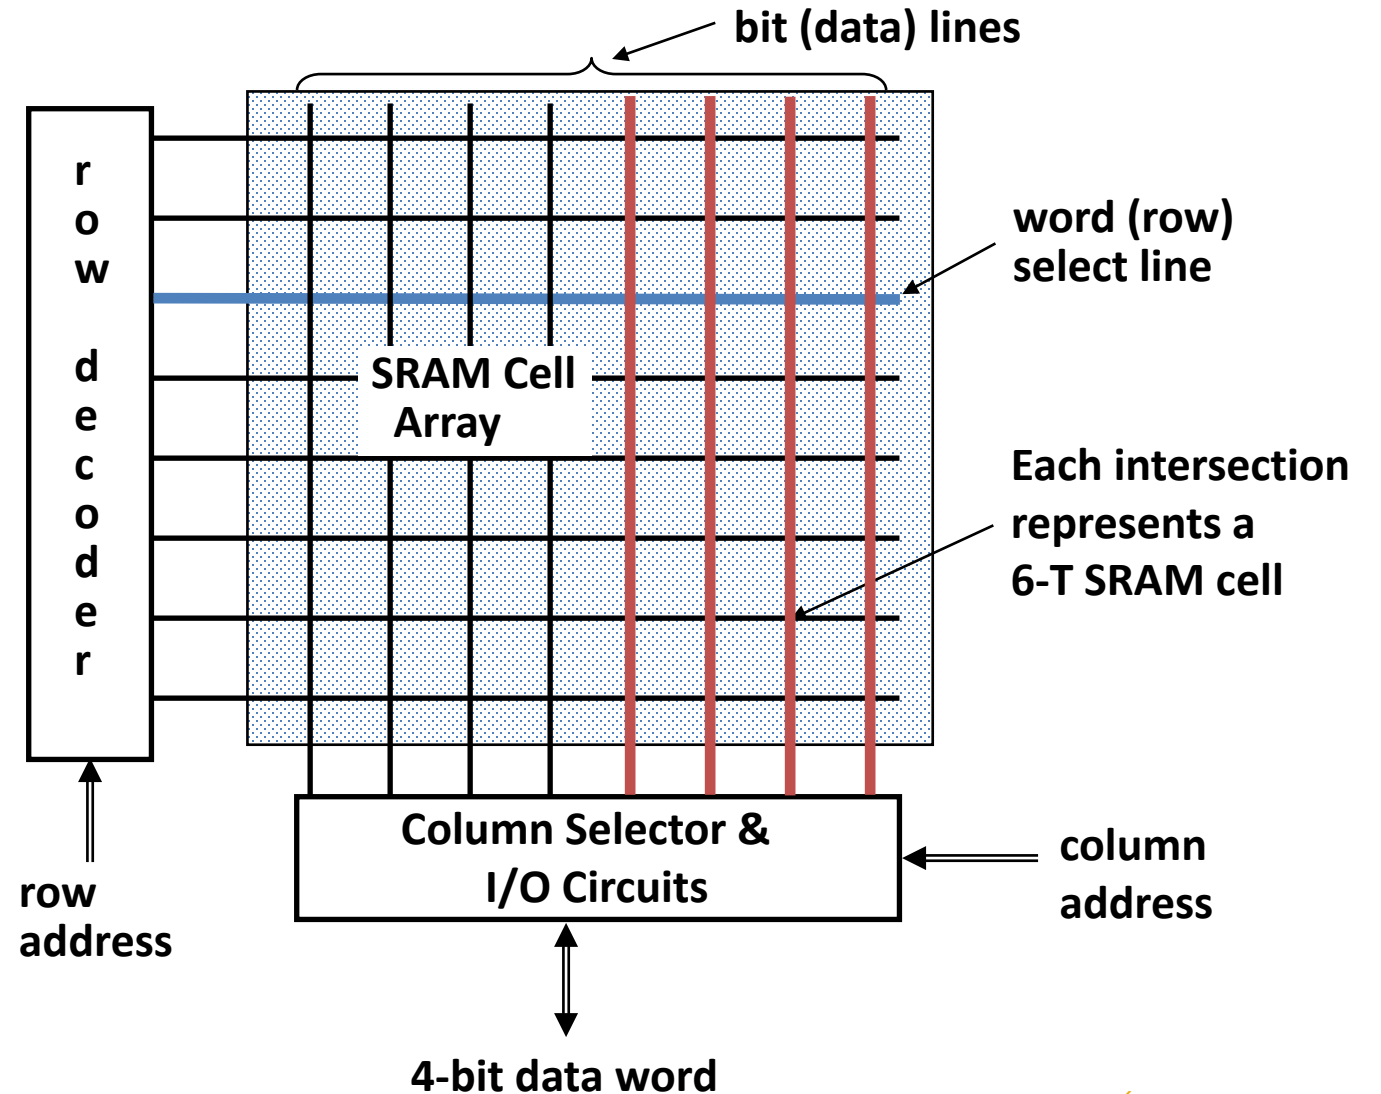
Cache Design

- **Caches use *SRAM* for speed and technology compatibility**
 - Low density (6 transistor cells), high power, expensive, fast
 - Static: content will last “forever” (until power turned off)
- **Example: A (2M X 16 bit) SRAM logic**



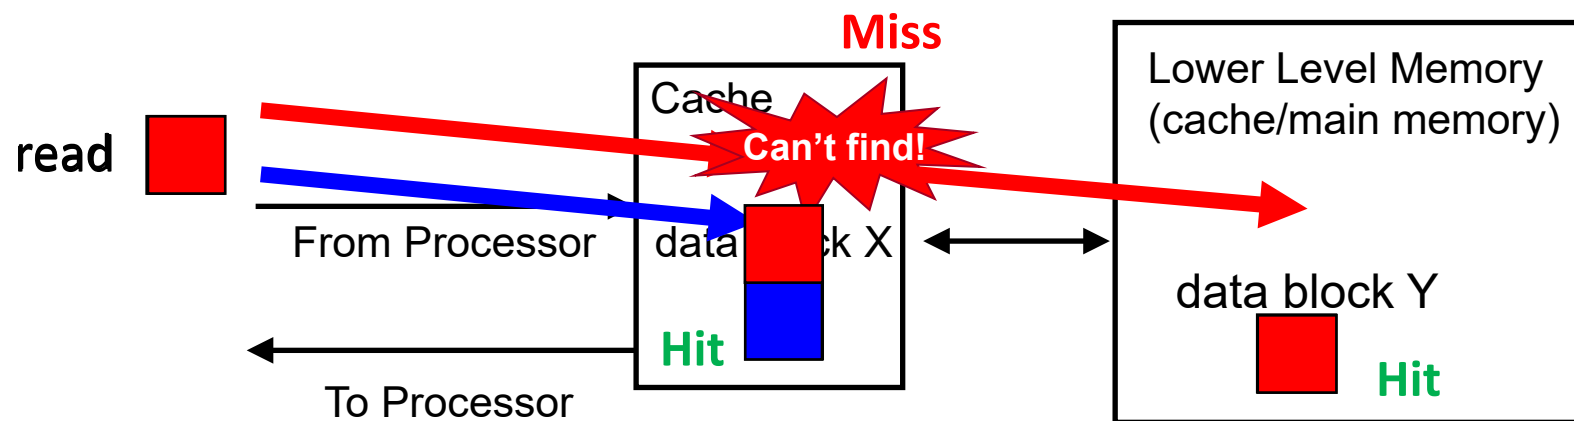
SRAM Cache Design

- Each row holds a data block
- Column address selects the requested word from block




Cache Hit and Miss

- **Hit:** Data appears in some block of the cache
 - **Hit Rate:** # hits / total accesses on the cache
 - **Hit Time:** Time to access the cache
- **Miss:** Data needs to be retrieved from the lower level (and stored in cache)
 - **Miss Rate:** 1 - (Hit Rate)
 - **Miss Penalty:** Average delay in the processor caused by each miss




Principle of Locality

- **Temporal Locality:** If an address is referenced, it tends to be referenced again
- **Spatial Locality:** If an address is referenced, neighboring addresses tend to be referenced
- **How to create a memory system that gives the illusion of being large, cheap and fast?**
 - With hierarchy
 - With parallelism




SRAM Cache Design

- Each row holds a data block
- Column address selects the requested word from block

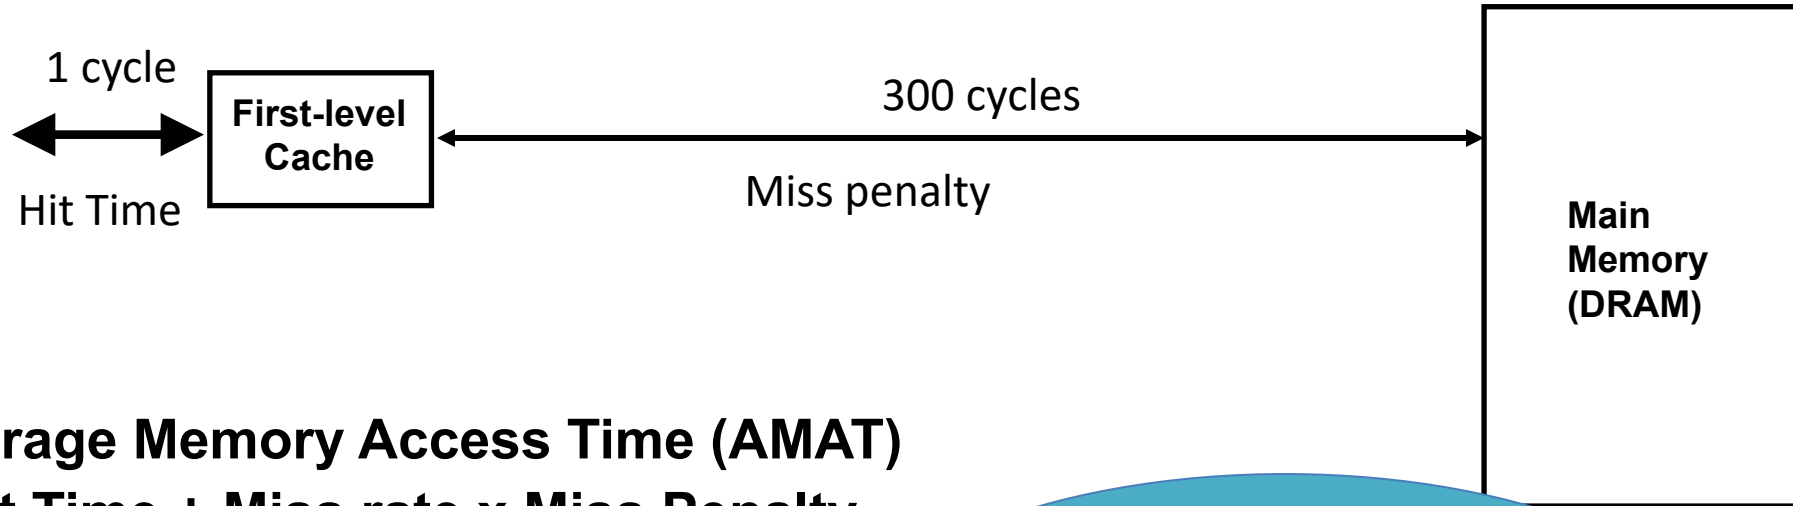


Cache Hit and Miss

- **Hit:** Data appears in some block of the cache
 - **Hit Rate:** # hits / total accesses on the cache
 - **Hit Time:** Time to access the cache
- **Miss:** Data needs to be retrieved from the lower level (and stored in cache)
 - **Miss Rate:** 1 - (Hit Rate)
 - **Miss Penalty:** Average delay in the processor caused by each miss



Memory Hierarchy Performance



- **Average Memory Access Time (AMAT)**
= Hit Time + Miss rate x Miss Penalty

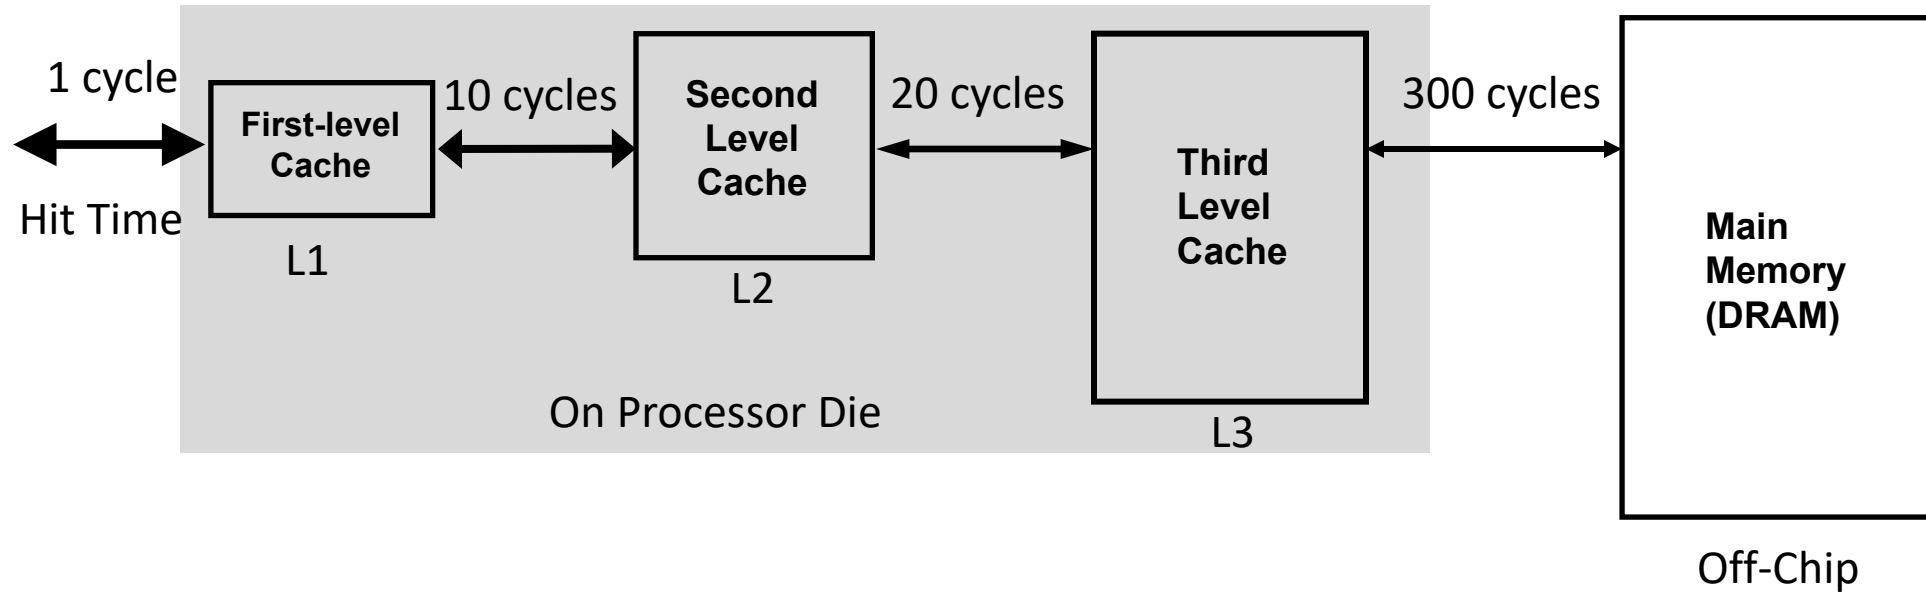
- **Example:**

- Cache Hit = 1 cycle
- Miss rate = 10% = 0.1
- Miss penalty = 300 cycles
- $AMAT = T_{hit}(L1) + Miss_rate(L1) \times T(Memory) = 1 + 0.1 \times 300 = 31$ cycles

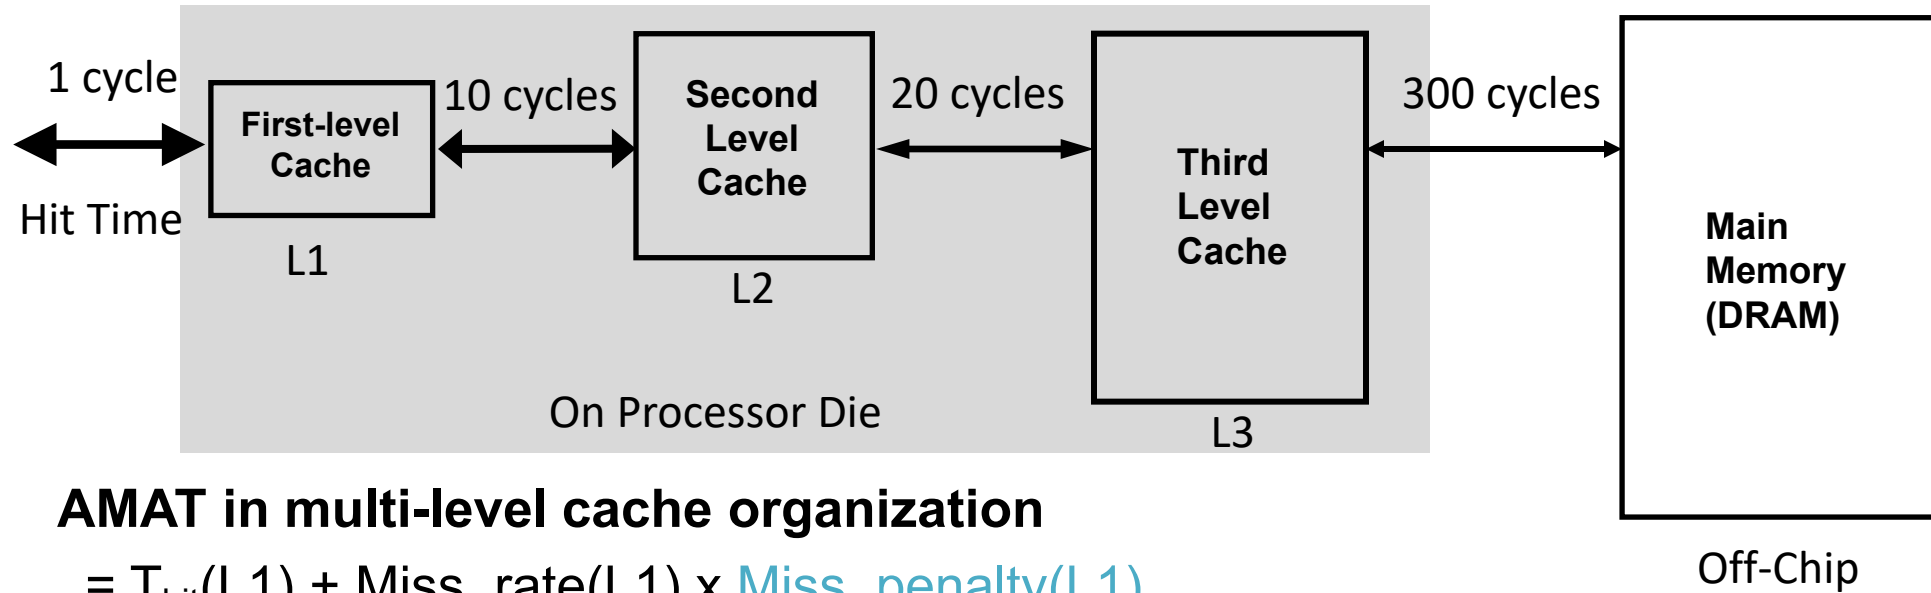
Due to long-latency memory, AMAT is 30 cycles longer than cache latency.

Can we reduce the overhead?

Reducing Penalty: Multi-Level Cache

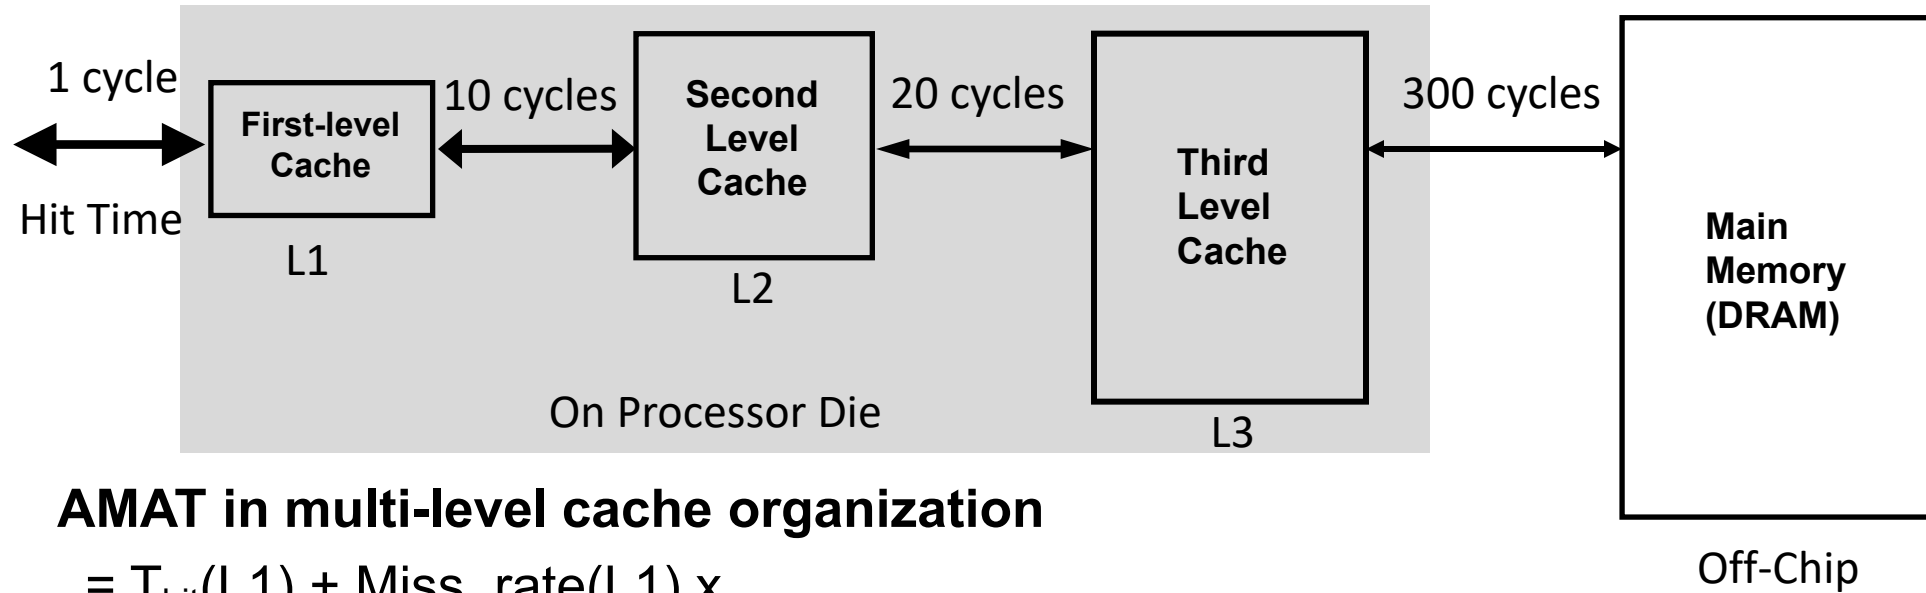


Reducing Penalty: Multi-Level Cache




- **AMAT in multi-level cache organization**
$$= T_{\text{hit}}(L1) + \text{Miss_rate}(L1) \times \text{Miss_penalty}(L1)$$

Reducing Penalty: Multi-Level Cache



- **AMAT in multi-level cache organization**
$$= T_{\text{hit}}(L1) + \text{Miss_rate}(L1) \times [T_{\text{hit}}(L2) + \text{Miss_rate}(L2) \times \text{Miss_penalty}(L2)]$$


Reducing Penalty: Multi-Level Cache



- **AMAT in multi-level cache organization**


$$\begin{aligned} &= T_{\text{hit}}(\text{L1}) + \text{Miss_rate}(\text{L1}) \times \\ &\quad [T_{\text{hit}}(\text{L2}) + \text{Miss_rate}(\text{L2}) \times \\ &\quad \{ T_{\text{hit}}(\text{L3}) + \text{Miss_rate}(\text{L3}) \times \text{Miss_penalty}(\text{L3}) \}] \end{aligned}$$

Reducing Penalty: Multi-Level Cache



- **AMAT in multi-level cache organization**
$$= T_{hit}(L1) + Miss_rate(L1) \times [T_{hit}(L2) + Miss_rate(L2) \times \{ T_{hit}(L3) + Miss_rate(L3) \times T(\text{memory}) \}]$$

Reducing Penalty: Multi-Level Cache



- **AMAT in multi-level cache organization**
$$= T_{hit}(L1) + Miss_rate(L1) \times [T_{hit}(L2) + Miss_rate(L2) \times \{ T_{hit}(L3) + Miss_rate(L3) \times T(memory) \}]$$

- **Example:**

- Miss rate of L1, L2, L3 = 10%, 5%, 1%, respectively
- $AMAT = 1 + 0.1 \times [10 + 0.05 \times \{ 20 + 0.01 \times 300 \}] = 2.115$ cycles

Vs. 31 cycles
14.7x speedup!

What to Keep in Caches?


- **It depends on the cache organization and replacing policy**
- **Cache organization:**
 - **Cache line (block):** The basic unit of data replacement. A longer cache line fetches and replaces more data per miss.
 - **Set:** An entry that one cache line is mapped to according to certain bits of its address.
 - **Way:** A slot within a cache set to hold one history cache line.
- **Cache line replacing policy:** which history line should be replaced?
 - In a set entry, each cache line can be identified with a **Tag** (a portion of its address).
 - Least recently used (LRU), First-in first-out (FIFO), Random, etc.

Cache Indexing

Example: a simple 2-set x 2-way cache

	Way 0	Way 1
Set 0	block 0	block 2
Set 1	block 1	block 3


Address decoding:



Cache Types

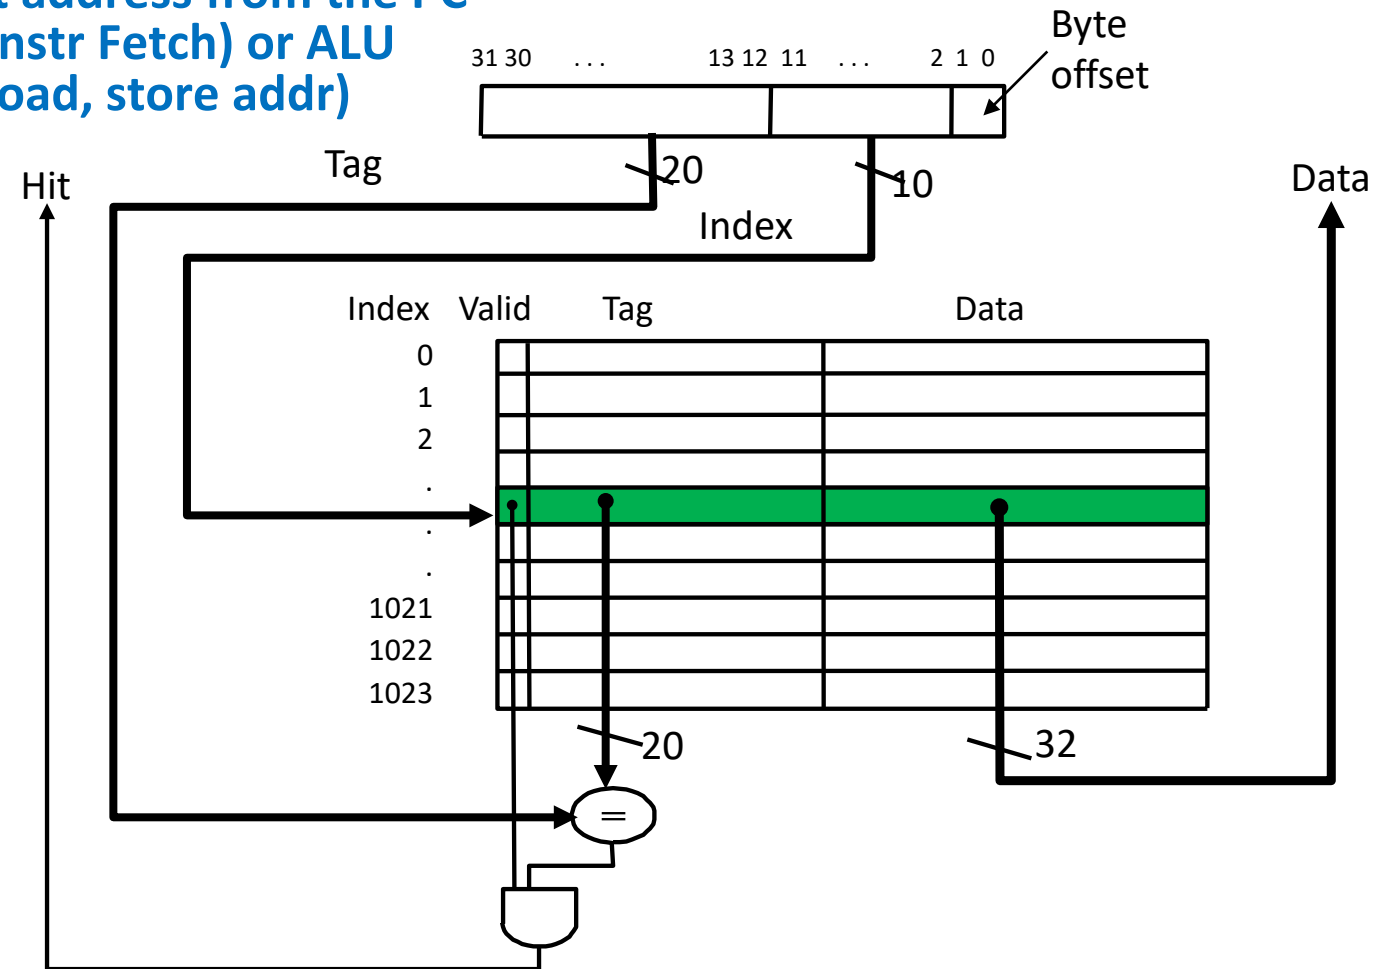
- **N-way Set-Associative:** Number of ways > 1 & Number of sets > 1
 - Slightly complex searching mechanism
- **Direct Mapped:** Number of ways = 1
 - Fast indexing mechanism
- **Fully-Associative:** Number of sets = 1
 - Extensive hardware resources required to search

	Way 0	Way 1	...
Set 0	block 0	block 2	
Set 1	block 1	block 3	
⋮			

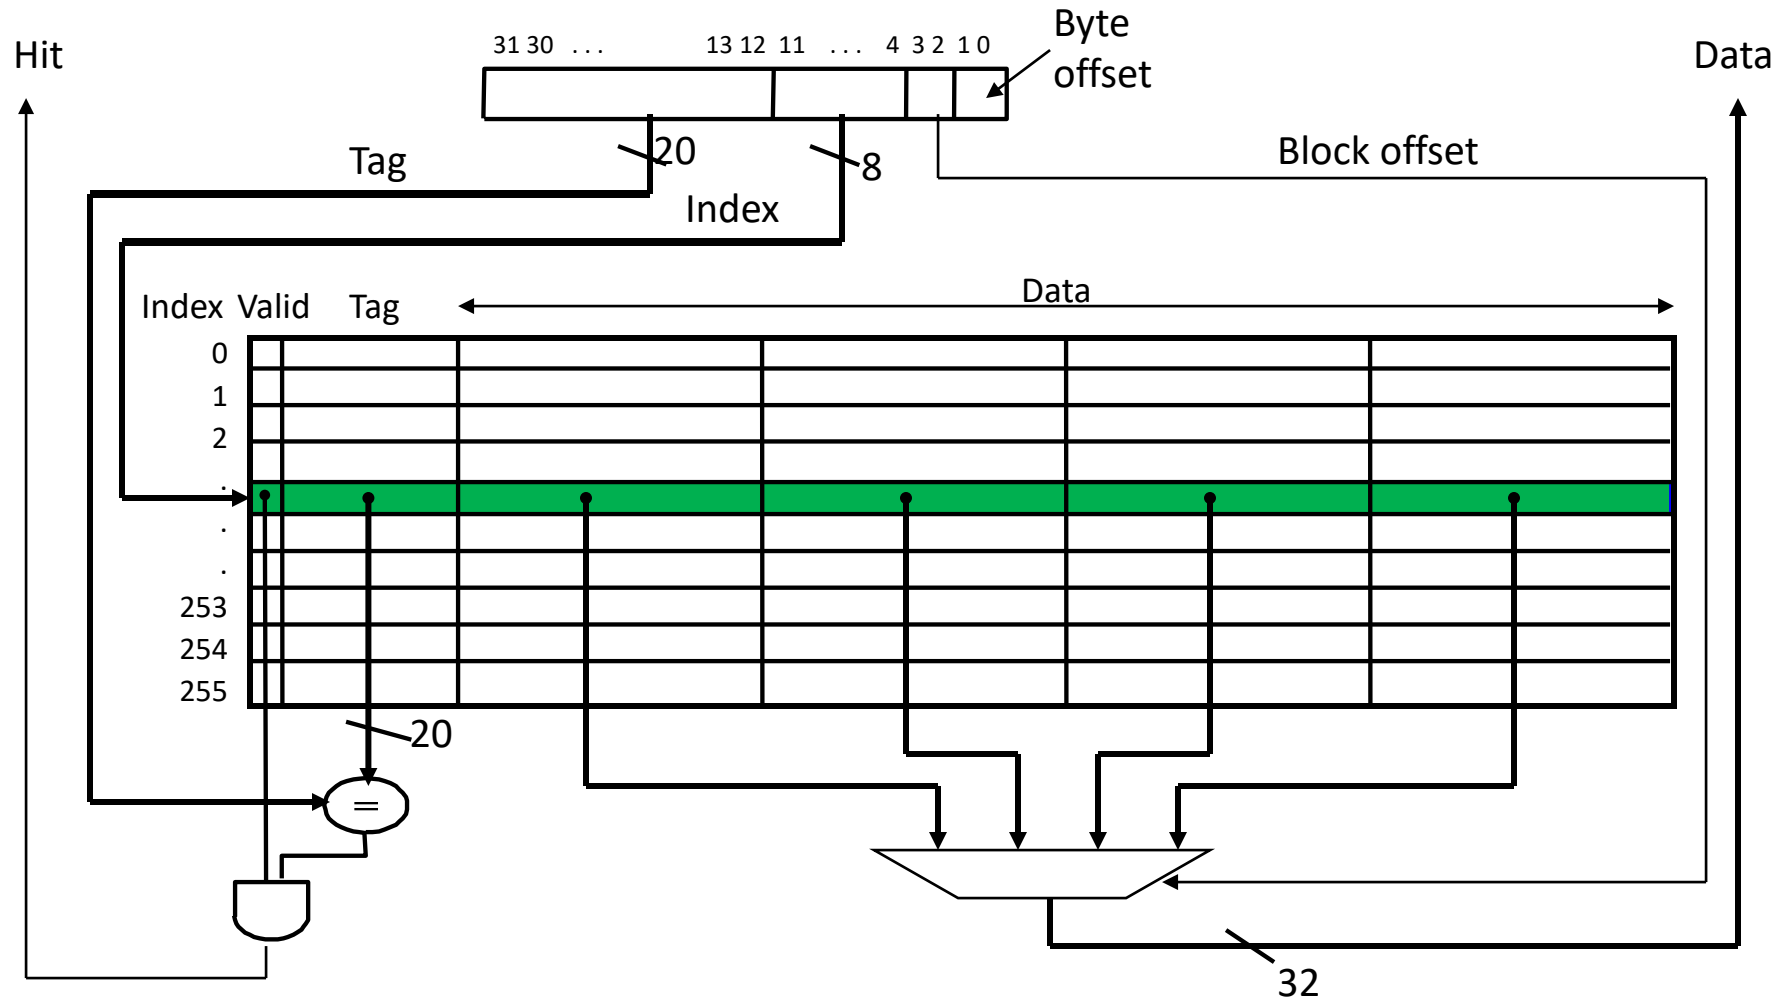


MIPS Direct Mapped Cache Example

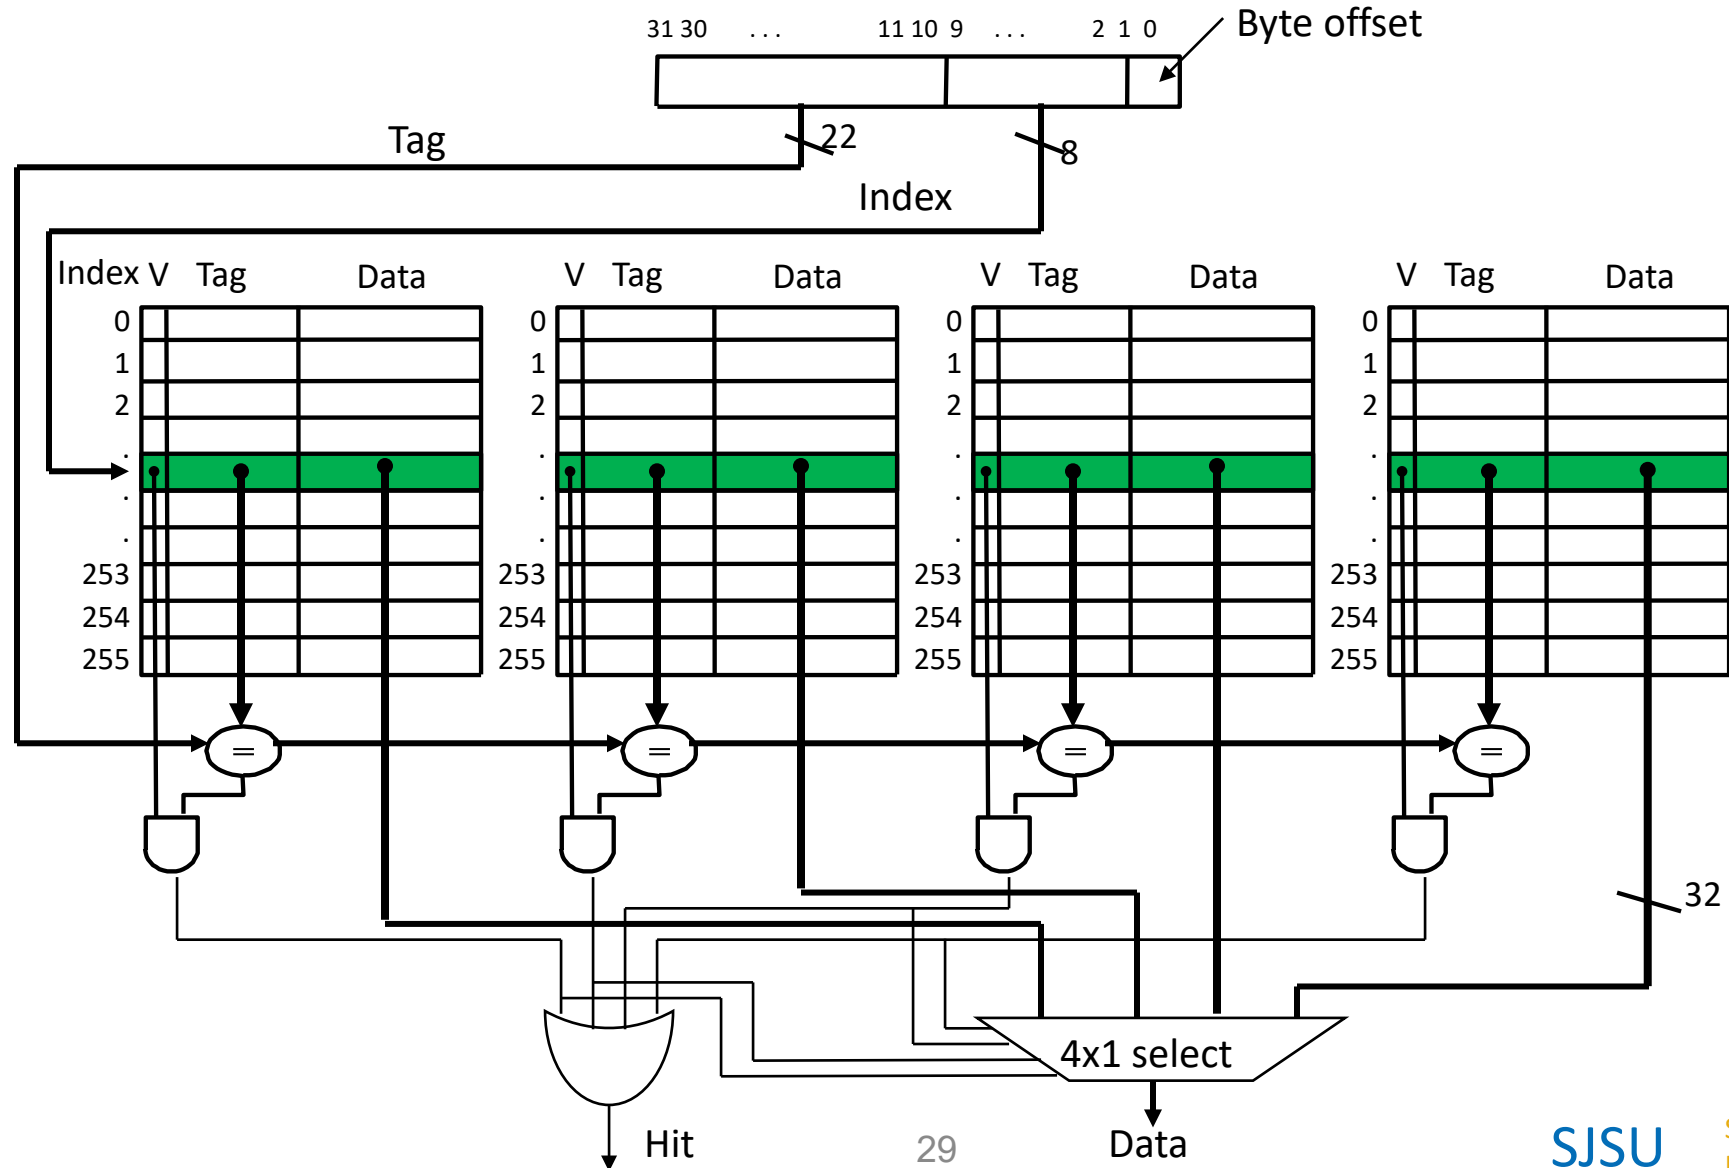
32bit address from the PC
(Instr Fetch) or ALU
(load, store addr)



Multiword Block Direct Mapped Cache



Four-Way Set Associative Cache



Costs of Set Associative Caches


- **Must have hardware to keep track of when each way's block was used relative to the other blocks in the set (for replacement policy)**
 - E.g., for 2-way set associative, takes **one bit per set** → set the bit when a block is referenced (and reset the other way's bit)
- **N-way set associative cache costs**
 - N comparators (delay and area)
 - MUX delay (set selection) before data is available
 - Data is available **after** set selection (and Hit/Miss decision). In a direct mapped cache, the cache block is available **before** the Hit/Miss decision
 - So its not possible to just assume a hit and continue and recover later if it was a miss
- **Total cache line size = valid field size + tag size + block data size + data for cache policy (e.g., time stamp, modified bit, etc.)**

Cache Miss Classification: The 3 C's


- **Compulsory (cold) Misses**
 - On the 1st reference to a block
 - Related to # blocks accessed by a code, not related to the configuration of a cache
- **Capacity Misses**
 - The program's working set size exceeds the cache capacity
- **Conflict Misses**
 - Multiple memory blocks map to the same set in set-associative caches

More Detailed Mapping Example

- We have a cache that has following configuration
 - Consists of 8 cache blocks (lines)
 - A data word is 4-byte
 - A block is 32-byte (8 words per block)
 - **Direct mapped**
- We load a word from 0x77FF1C68




- 1st step: Find data block no. that the word belongs to

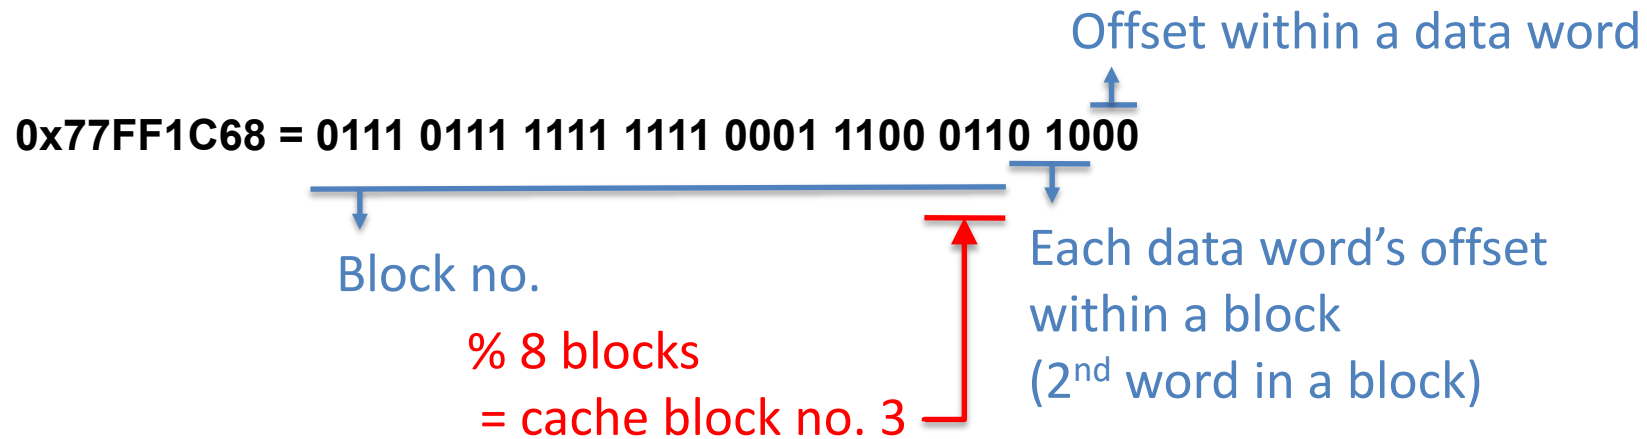


More Detailed Mapping Example

- We have a cache that has following configuration
 - Consists of 8 cache blocks (lines)
 - A data word is 4-byte
 - A block is 32-byte (8 words per block)
 - **Direct mapped**
- We load a word from 0x77FF1C68




- 2nd step: Map the block no. to cache block no.



More Detailed Mapping Example


- We have a cache that has following configuration
 - Consists of 8 cache blocks (lines)
 - A data word is 4-byte
 - A block is 32-byte (8 words per block)
 - **4-way Set-Associative**
- We load a word from 0x77FF1C68




- 1st step: Calculate how many sets exist
 $8 \text{ blocks} / 4 \text{ entries per set} = 2 \text{ sets}$

More Detailed Mapping Example

- We have a cache that has following configuration
 - Consists of 8 cache blocks (lines)
 - A data word is 4-byte
 - A block is 32-byte (8 words per block)
 - **4-way Set-Associative**
- We load a word from 0x77FF1C68



- 2nd step: Find data block no. that the word belongs to




More Detailed Mapping Example

- We have a cache that has following configuration
 - Consists of 8 cache blocks (lines)
 - A data word is 4-byte
 - A block is 32-byte (8 words per block)
 - **4-way Set-Associative**
- We load a word from 0x77FF1C68


pick any
available block
within set 1

cache blocks




set no. 0 1

- 3rd step: Map the block no. to cache block no.



More Detailed Mapping Example

- We have a cache that has following configuration
 - Consists of 8 cache blocks (lines)
 - A data word is 4-byte
 - A block is 32-byte (8 words per block)
 - **Fully-Associative**
- We load a word from 0x77FF1C68



- 1st and all step: Find any available cache block and map

Cache Replacement Policy

- **When loading a new block (on miss), if the cache is already full, which block should be replaced (in the set)?**
 - Random: Replace a randomly chosen line
 - FIFO: Replace the oldest line
 - LRU (Least Recently Used): Replace the least recently used line
 - Many others: Round Robin, LIFO, MRU, etc..

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A								
	1									
	2									
	3 (LRU)									
cache miss		m								

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B							
	1		A							
	2									
	3 (LRU)									
cache miss		m	m							

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C						
	1		A	B						
	2			A						
	3 (LRU)									
cache miss		m	m	m						

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C	D					
	1		A	B	C					
	2			A	B					
	3 (LRU)				A					
cache miss		m	m	m	m					

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C	D	C				
	1		A	B	C	D				
	2			A	B	B				
	3 (LRU)				A	A				
cache miss		m	m	m	m	h				

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C	D	C	D			
	1		A	B	C	D	C			
	2			A	B	B	B			
	3 (LRU)				A	A	A			
cache miss		m	m	m	m	h	h			

LRU to be replaced for E

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C	D	C	D	E		
	1		A	B	C	D	C	D		
	2			A	B	B	B	C		
	3 (LRU)				A	A	A	B		
cache miss		m	m	m	m	h	h	m		

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C	D	C	D	E	C	
	1		A	B	C	D	C	D	E	
	2			A	B	B	B	C	D	
	3 (LRU)				A	A	A	B	B	
cache miss		m	m	m	m	h	h	m	h	

LRU to be replaced for G

LRU Example

- **Cache configuration:**
 - Size: 4 blocks
 - Mapping: fully-associative
 - Replacement: LRU
- **Mem access sequence: (each character indicates a block address)**
 - ABCDCDECG

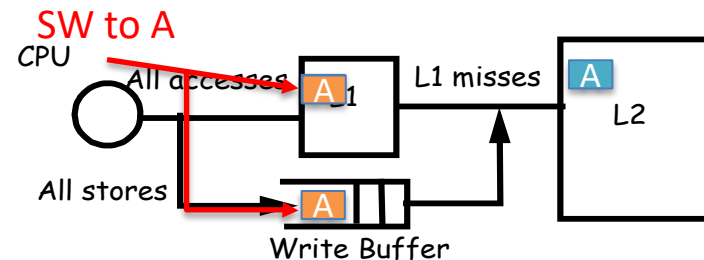
Accesses		A	B	C	D	C	D	E	C	G
priority order (LRU)	0 (MRU)	A	B	C	D	C	D	E	C	G
	1		A	B	C	D	C	D	E	C
	2			A	B	B	B	C	D	E
	3 (LRU)				A	A	A	B	B	D
cache miss		m	m	m	m	h	h	m	h	m

Cache Policies

- **Read policies:**
 - Read Hit: this is what we want
 - Read Miss: replacement policy
- **Write policies (L1D only): consistency issues & performance tradeoffs**
 - Write-through vs. write-back vs. write-evict (no-write)
 - also slightly different between write hit and write miss
 - Write Hit: write policy
 - Write Miss: write allocation policies (2 policies to categorize 3 write policies) + replacement policy


Cache Write Policy

- **Write-through: consistent with lower level**
 - The value is written to both the cache line and the lower-level memory.
 - No need to write a full block
 - Write buffer can be used so cache does not need to wait for the write to lower level
 - stall only if the write buffer is full
 - Example: assume cache block containing word address A is initially in the cache



Cache Write Policy

- **Write-through: consistent with lower level**
 - The value is written to both the cache line and the lower-level memory.
 - No need to write a full block
 - Write buffer can be used so cache does not need to wait for the write to lower level
 - stall only if the write buffer is full
 - Example: assume cache block containing word address A is initially in the cache



Cache Write Policy

- **Write-back: inconsistent with lower level**

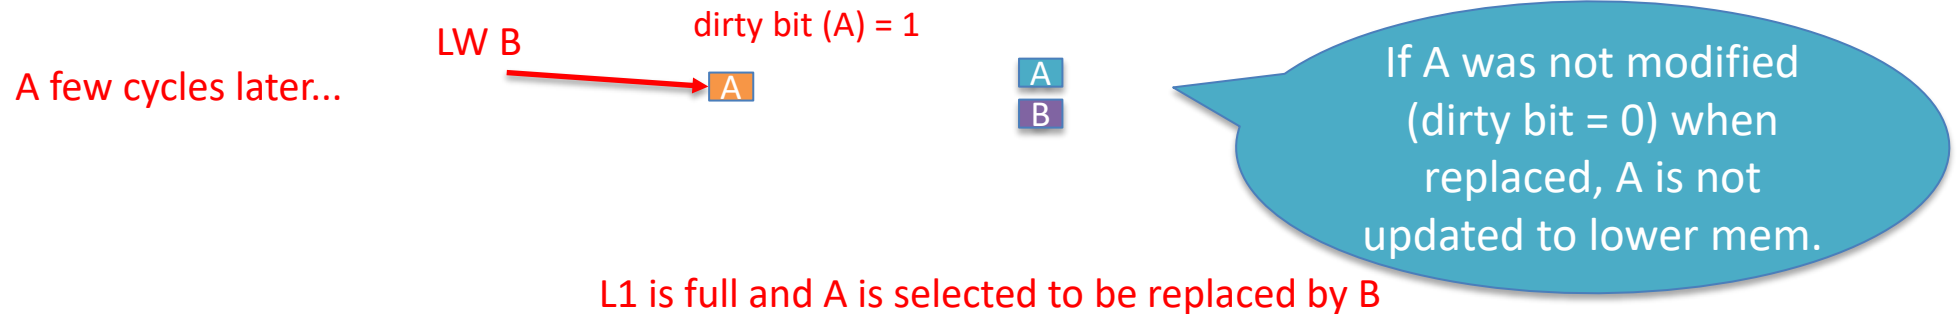
- The value is written only to the cache line. The modified (dirty) cache line is written to main memory only when it is evicted.
 - 1 dirty bit is needed for each cache line.
- A write-back buffer to update lower level with evicted dirty blocks
 - Must write a full block at this point since we do not know which word is modified
- Example: assume cache block containing word address A is initially in the cache



Cache Write Policy

- **Write-back: inconsistent with lower level**

- The value is written only to the cache line. The modified (dirty) cache line is written to main memory only when it is evicted.
 - 1 dirty bit is needed for each cache line.
- A write-back buffer to update lower level with evicted dirty blocks
 - Must write a full block at this point since we do not know which word is modified
- Example: assume cache block containing word address A is initially in the cache



Cache Write Policy

- **Write-evict (no-write/bypassing): no consistency issue**
 - The value is written only to the lower level.
 - If the block containing the write address exists in the cache, evict it.
 - The write only need to send the data specified by the store instruction.
 - A write buffer can also be used to reduce the delay.

Cache Policies

- **Read policies:**
 - Read Hit: this is what we want
 - Read Miss: replacement policy + write policy of evicted data
- **Write policies: consistency issues & performance tradeoffs**
 - Write-through vs. write-back vs. write-evict (no-write)
 - also slightly different between write hit and write miss
 - Write Hit: write policy
 - Write Miss: write allocation policy (+ replacement policy) + write policy of evicted data

Cache Write Policy

- **Read miss:**
 - Write-through: find victim block + fetch block from lower level
 - write-back: find victim block + write back evicted block if dirty + fetch block from lower level
 - No-write: find victim block + fetch block from lower level
- **Write miss: write allocation policies**
 - Write-allocate:
 - Write-through: find victim block + fetch block from lower level + store value to block and lower level
 - write-back: find victim block + write back evicted block if dirty + fetch block from lower level + store value to block
 - No-write-allocate: bypass the cache, store value to lower level directly
- **Write hit and write miss can use different policies**
 - e.g., write-hit write-evict + write-miss write-allocate-write-back

Reduce Miss Rate (1): Code Optimization

- **Misses occur if sequentially accessed array elements come from different cache blocks**
- **Code optimizations → No hardware change**
 - Rely on programmers or compilers
- **Examples:**
 - Loop interchange: In nested loops, outer loop becomes inner loop and vice versa
 - Loop blocking: partition large array into smaller blocks, thus fitting the accessed array elements into cache size

Loop Interchange

Assume: cache line size = 5 words

What addresses are accessed in each iteration of inner loop?



```
/* Before */  
for (j=0; j<5; j++)  
  for (i=0; i<5; i++)  
    x[i][j] = 2*x[i][j]
```

Column-major ordering



Row-major ordering


```
/* After */  
for (i=0; i<5; i++)  
  for (j=0; j<5; j++)  
    x[i][j] = 2*x[i][j]
```



Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```

The second matrix may always encounter misses because individual data are from different cache blocks




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the traditional matrix multiplication:

```
/* Before */  
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    for (k=0; k<N; k++)  
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```

If N is large, the cache line that has been accessed is likely to be replaced before next access




To generate one output element (red color), the 6 elements of each input matrix (gray color) will be accessed.

Loop Blocking

In the block-based matrix multiplication:

```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```


Matrix is computed in
a smaller block unit
→ higher locality



Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



Loop Blocking

In the block-based matrix multiplication:


```
/* After (block size = 4) */  
for (i=0; i<N; i=i+b)  
  for (j=0; j<N; j=j+b)  
    for (k=0; k<N; k=k+b)  
      for (ii=i; ii<i+b; ii++)  
        for (jj=j; jj<j+b; jj++)  
          for (kk=k; kk<k+b; kk++)  
            C[ii][jj] = C[ii][jj] + A[ii][kk] * B[kk][jj];
```



Reduce Miss Rate (2): Reduce the 3 C's


- **Increase Cache Size**
 - Reduce miss for: capacity miss, conflict miss
 - But has many limitations
- **Increase Associativity**
 - Reduce miss for: conflict miss
 - But may increase access latency
- **Increase Cache Block Size**
 - Reduce miss for: compulsory
 - But may increase miss penalty (more data will be evicted and fetched)
 - Very large blocks could increase miss rate

Miss Rate vs Block Size vs Cache Size




- Cache size – the larger the better
- Block size – tradeoffs

Block Size Tradeoff



$$\text{Average Memory Access Time} = \text{Hit Time} + \text{Miss Penalty} \times \text{Miss Rate}$$

Benefits of Set Associative Caches



Largest gains when going from direct mapped to 2-way (20%+ reduction in miss rate)

Reduce Miss Rate (3): Policies

- **Use appropriate policies**
 - LRU + write-back works for most applications
 - Many GPUs use write-hit write-evict + write-miss no-write-allocate
 - Can even dynamically change policy according to profiling
- **Use more advanced policies**
 - E.g., ML based

Reduce Miss Rate (4): Multi-level Caches

- Having a unified L2 cache (i.e., it holds both instructions and data) and in some cases even a unified L3 cache
- L1 cache should focus on **minimizing hit time** in support of a shorter clock cycle
- Secondary cache(s) should focus on **reducing miss rate** to reduce the penalty of long main memory access times

Split vs. Unified Caches

- **Split L1I\$ and L1D\$: to optimize t_{hit}**
 - Low capacity/associativity/block size (reduce t_{hit})
 - Minimize structural hazards
 - Can optimize L1I\$ for wide output and no writes
- **Unified L2: to optimize hit rate**
 - t_{hit} is less important due to less frequent accesses & long delay already
 - High capacity/associativity/block size (reduce $\%_{miss}$)
 - Unused instr capacity can be used for data (fewer capacity misses)
 - Instr / data conflicts (small to no increase for conflict misses in a large cache)
 - Instr / data structural hazards are rare (would take a simultaneous L1I\$ and L1D\$ miss)

Reduce Miss Rate (5): Prefetching

- **Hardware prefetching**
 - Fetch blocks into the cache proactively (speculatively)
 - Key is to anticipate the upcoming miss addresses accurately
 - Relies on having unused memory bandwidth available
- **A simple case is to use **next block** prefetching**
 - Miss on address $X \rightarrow$ anticipate next reference miss on $X + \text{block-size}$
 - Works well for instr's (sequential execution) and for arrays of data
- **Need to initiate prefetches sufficiently in advance**
- **If prefetched instr/data is not used, the cache is polluted with unnecessary data (possibly evicting useful data)**

Measuring Performance with Caches

- Assuming cache hit costs are included as part of the normal CPU execution cycle, then

$$\begin{aligned}\text{CPU time} &= \text{IC} \times \text{CPI} \times \text{CP} \\ &= \text{IC} \times (\underbrace{\text{CPI}_{\text{ideal}} + \text{Memory-stall cycles}}_{\text{CPI}_{\text{stall}}}) \times \text{CP}\end{aligned}$$

Memory-stall cycles come from cache misses (read-stalls and write-stalls)

Note: this is miss ratio with regard to all instructions
= read ratio * cache miss rate

$$\text{Read-stall cycles} = \text{read miss ratio} \times \text{read miss penalty}$$

$$\text{Write-stall cycles} = \text{write miss ratio} \times \text{write miss penalty} + \text{write buffer stalls}$$

For write-through caches, we can simplify this to

$$\text{Memory-stall cycles} = \text{miss ratio} \times \text{miss penalty}$$

Impacts of Cache Performance

- **Relative cache penalty increases as processor performance improves (faster clock rate and/or lower CPI)**
 - The memory speed is unlikely to improve as fast as processor cycle time. When calculating CPI_{stall} , the cache miss penalty is measured in *processor* clock cycles needed to handle a miss
 - The lower the CPI_{ideal} , the higher the impact of stalls
- **Example: A processor with a CPI_{ideal} of 2, a 100 cycle miss penalty, 36% load/store instructions, and 2% I\$ and 4% D\$ miss rates**

Memory-stall cycles = $2\% \times 100 + 36\% \times 4\% \times 100 = 3.44$

$CPI_{stalls} = 2 + 3.44 = 5.44$
- **What if the CPI_{ideal} is reduced to 1? Or the processor clock rate is doubled (doubling the miss penalty)?**


Other than Miss Rates

- **Cache bypassing (e.g., L1D)**
 - Reduce latency if L1 misses
- **Sector cache**
 - Fetch only a portion of the cache line
 - Not the same as reducing cache line size
- **Value prediction**
 - Using history value to predict future values
 - Roll back if predict wrong
 - Or accept a certain quality loss if the application is error tolerant

DRAM Subsystem Organization

Dram: Dynamic RAM


DRAM Organization:



- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column

Connected to an on-die memory controller

DRAM Cells

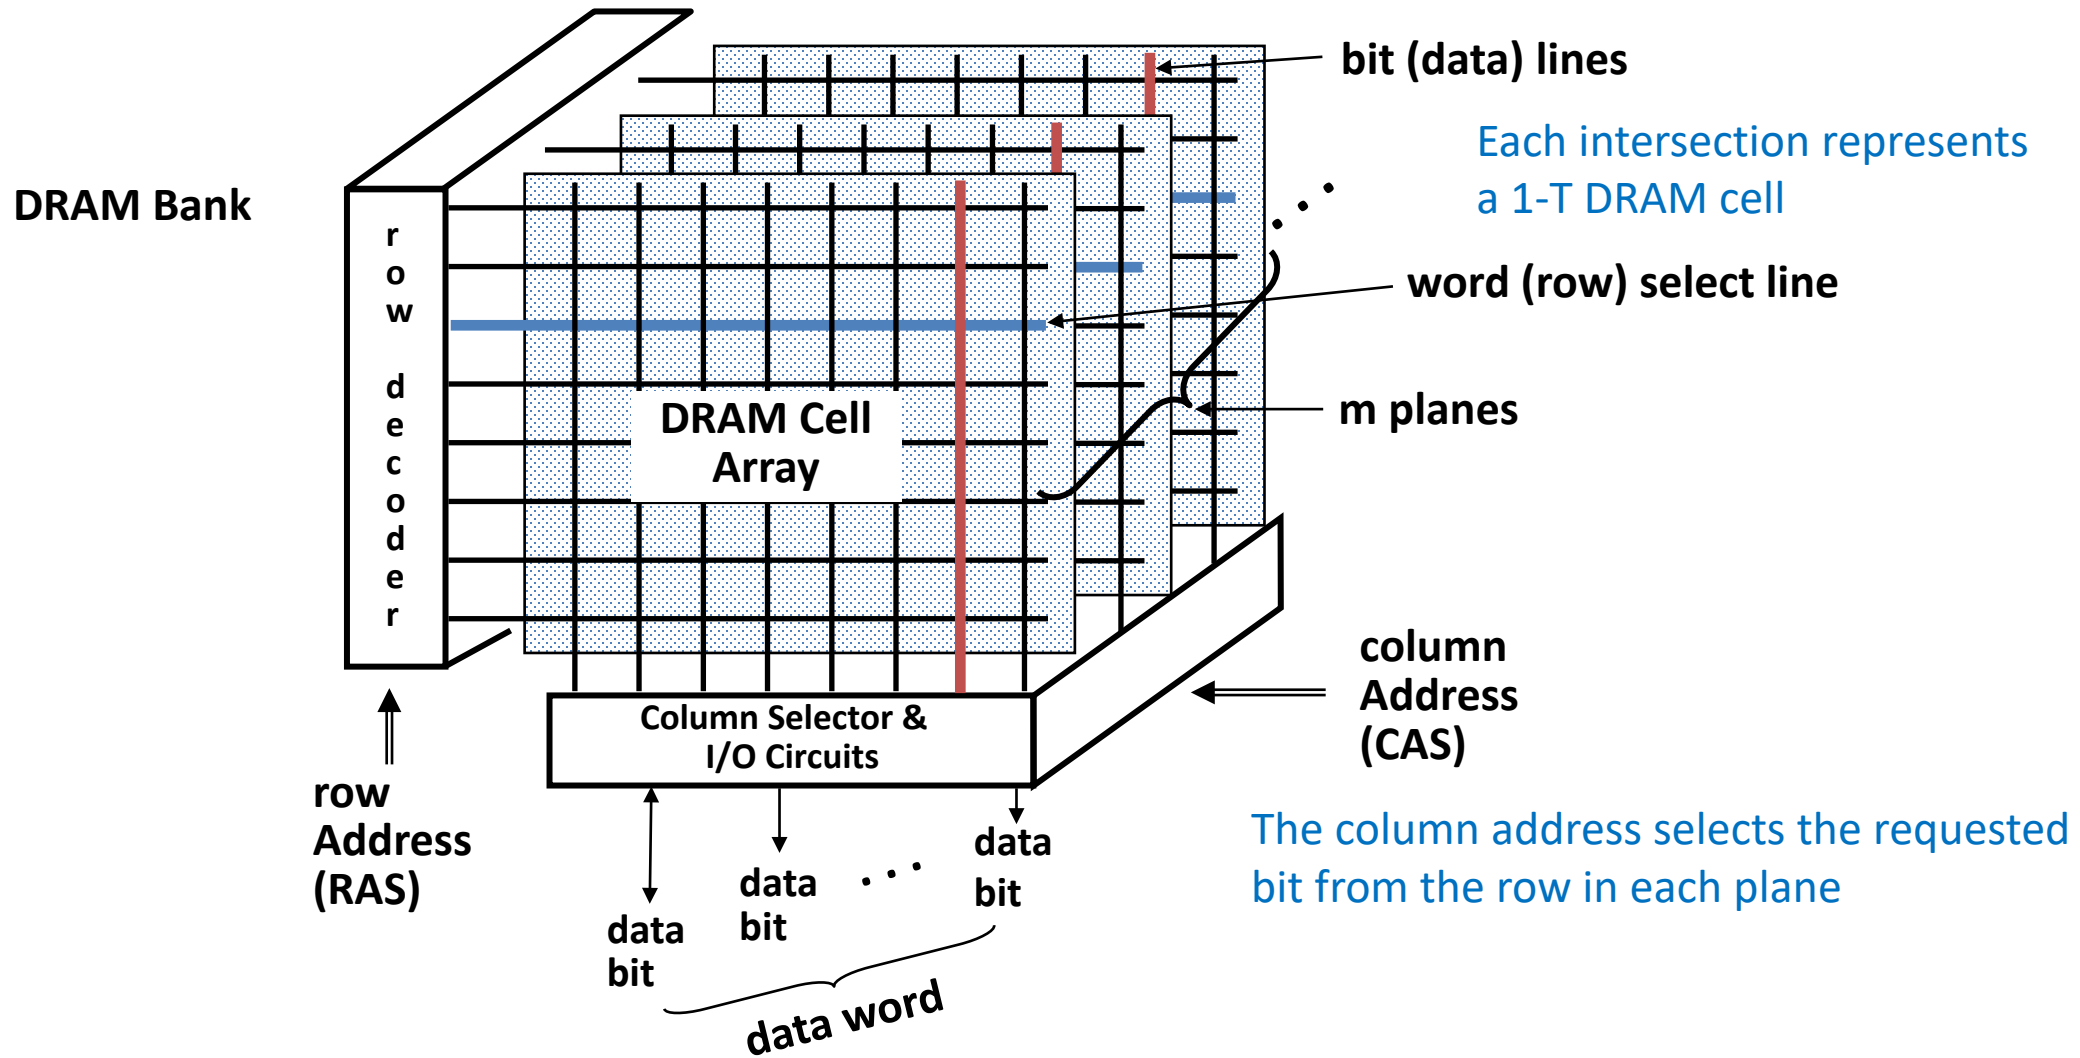


- **A DRAM cell consists of a capacitor and an access transistor**
- **It stores data in terms of charge in the capacitor**
- **Cheaper and larger than SRAM**
 - A DRAM chip consists of (10s of 1000s of) rows of such cells

DRAM Refresh

- **DRAM capacitor charge leaks over time**
- **The memory controller needs to refresh each row periodically to restore charge**
 - Dynamic (i.e., never in a stable state)
 - Activate each row every N ms (e.g., typical $N = 64$ ms)
- **Downsides of refresh**
 - **Energy consumption**: Each refresh consumes energy
 - **Performance degradation**: DRAM rank/bank unavailable while refreshed
 - **QoS/predictability impact**: (Long) pause times during refresh
 - **Refresh rate limits DRAM capacity scaling**

DRAM Design 1: The Classical



DRAM Performance Metrics

- **DRAM addresses are divided into 2 halves (row and column)**
 - *RAS* or *Row Access Strobe* that triggers the row decoder
 - *CAS* or *Column Access Strobe* that triggers the column selector
- **Latency: Time to access one word**
 - *Access Time*: time between request and when word is read or written
 - read access and write access times can be different
 - *Cycle Time*: time between successive (read or write) requests
 - Usually, cycle time > access time
- **Bandwidth: How much data can be supplied per unit time**
 - Width of the data channel * channel usage frequency


DRAM Performance Metrics

- **DRAM addresses are divided into 2 halves (row and column)**
 - *RAS* or *Row Access Strobe* that triggers the row decoder
 - *CAS* or *Column Access Strobe* that triggers the column selector
- **Latency: Time to access one word**
 - *Access Time*: time between request and when word is read or written
 - read access and write access times can be different
 - *Cycle Time*: time between successive (read or write) requests
 - Usually, cycle time > access time
- **Bandwidth: How much data can be supplied per unit time**
 - Width of the data channel * channel usage frequency

DRAM Design 1: The Classical

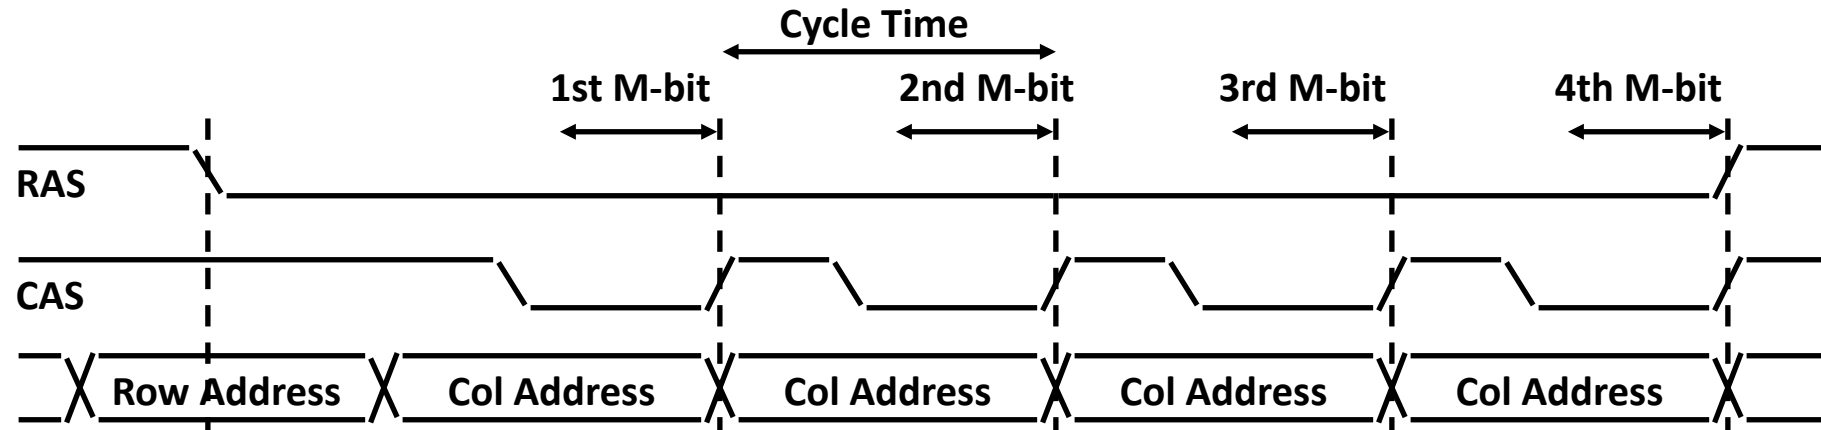
- **DRAM Organization:**

- N rows x N column x M-bit (planes)
- Reads or Writes M-bit at a time
- Each M-bit access requires a RAS / CAS cycle



DRAM Design 1: The Classical

- **Page Mode: A row is kept “open” by keeping the RAS asserted**
 - Pulse CAS to access other M-bit blocks on **that** row
 - Successive reads or writes within the row are faster since don't have to precharge and (re)access that row




DRAM Design 2: Synchronous DRAMs

- Like page mode DRAMs, synchronous DRAMs (SDRAMs) can transfer a **burst** of data from a series of sequential addresses in the **same** row
- For words in the same burst, don't have to provide the complete (row and column) addresses
 - Specify the starting (row+column) address and the burst length (burst must be in the same row). The row is accessed from the DRAM and loaded into a row buffer (SRAM).
 - Data words in the burst are then accessed from that SRAM under control of a clock signal.


Synchronous DRAM (SDRAM) Operation

- **After RAS loads a row into the SRAM cache**
 - Input CAS as the starting “burst” address along with a burst length to read a burst of data from a series of sequential addresses within that row on the clock edge



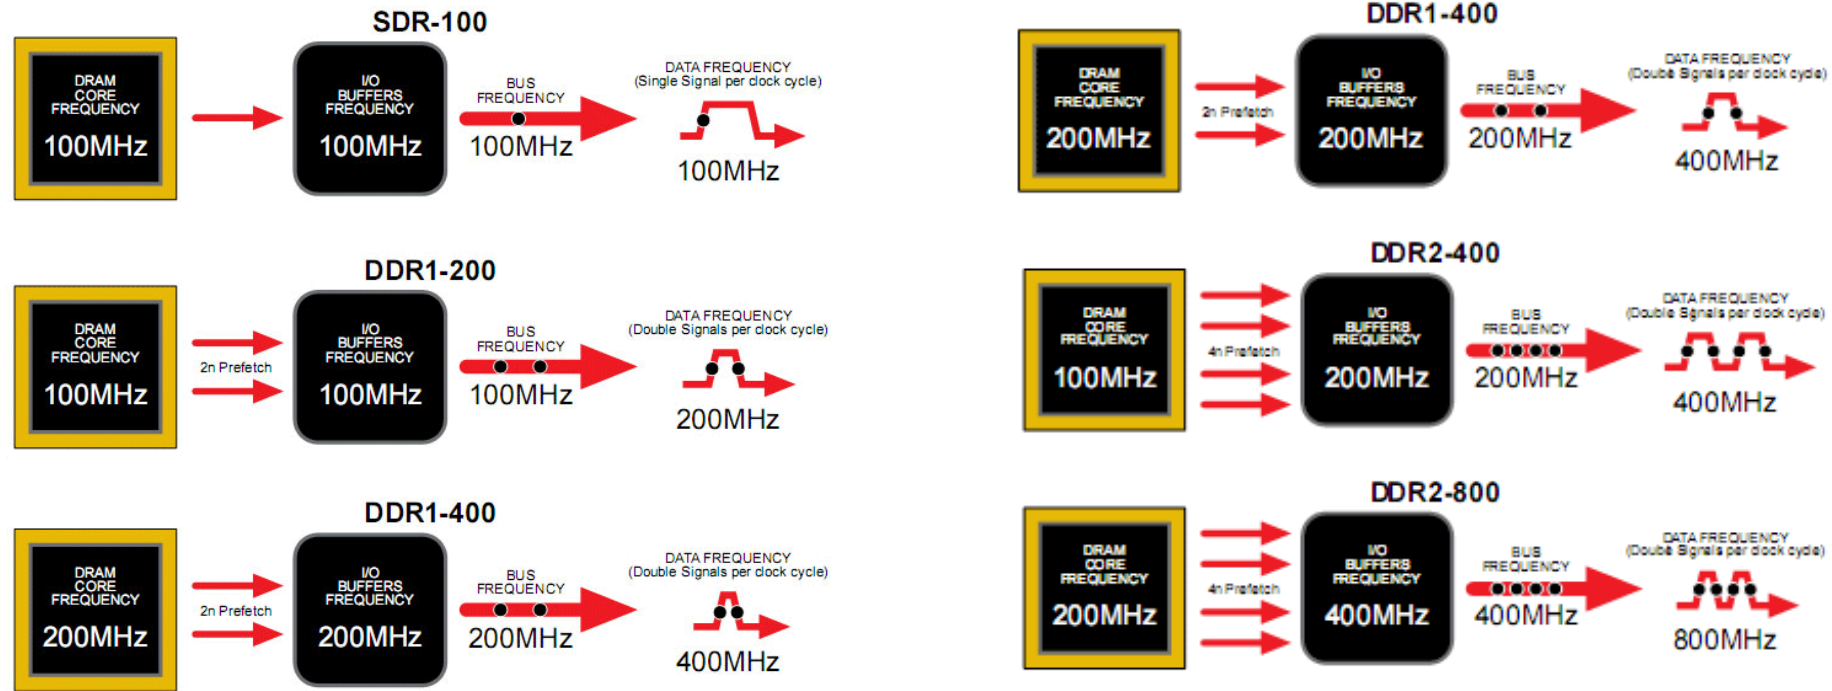
DDR (Double Data Rate) SDRAMs

- Transfers burst data on both the rising and falling edge of the clock (so twice fast)
 - 2n core prefetch



DDR (Double Data Rate) SDRAMs

- **DDR1 VS DDR1+:**



▲ Simplified Comparison between SDR-100, DDR1-200 and DDR1-400
Illustration: Ryan J. Leng


▲ Simplified Comparison between DDR1-400, DDR2-400 and DDR2-800
Illustration: Ryan J. Leng

- **DDR2 vs. QDR**

DRAM Subsystem Organization

Dram: Dynamic RAM

DRAM Organization:



- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column


- **Chip:**

- Consists of multiple banks (2-16 in Synchronous DRAM)
 - Banks work in parallel to overlap delay
- Banks share command/address/data buses
- The chip itself has a narrow interface (4-16 bits per read)

DRAM Subsystem Organization

Dram: Dynamic RAM

DRAM Organization:




- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column

- **Rank:**
 - Multiple chips operated together to form a wide interface
 - All chips comprising a rank are controlled at the same time
 - Respond to a single command
 - Share address and command buses, but provide different data
 - E.g., Using 8 chips with 8-bit interface to form a 64-bit wide Rank

DRAM Subsystem Organization

Dram: Dynamic RAM


DRAM Organization:



- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column

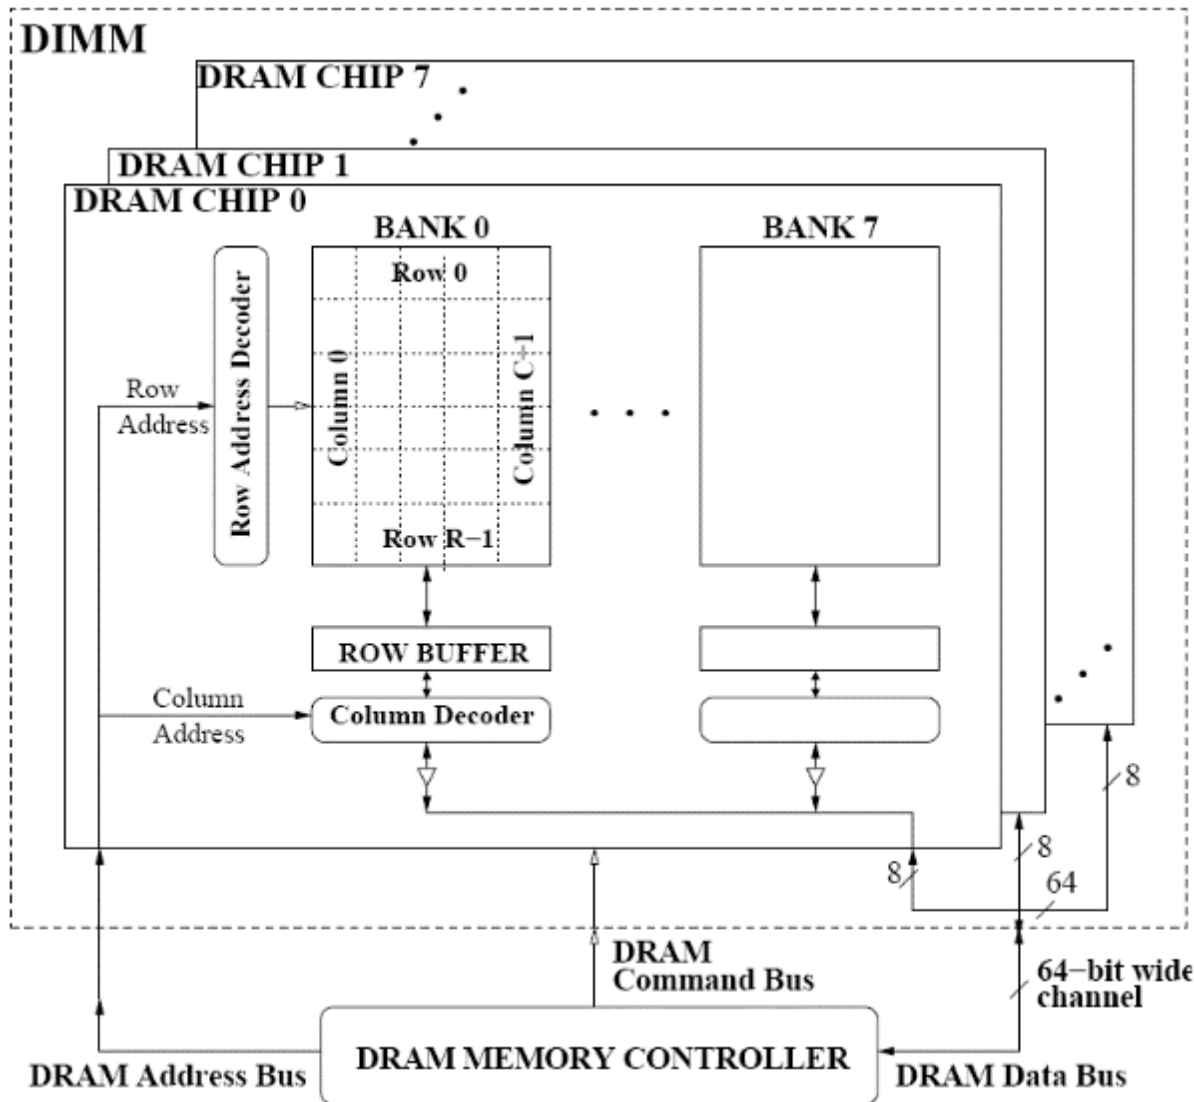
- **DIMM (dual inline memory module):**
 - A DRAM module consists of one or more ranks
 - This is what you plug into your motherboard

DIMM (Dual Inline Memory Module)



- **Contains DRAM chips each have data widths of x4 or x8**
 - Can be on both sides
- **Can have more than one rank**
 - Increase storage capacity
 - Only one rank accessible at a time
- **SIMM vs DIMM**
 - DIMM has separate contacts on each side of the board to provides twice as much data rate.

A 64-bit Wide DIMM (One Rank)



- **Advantages:**

- Acts like a **high-capacity DRAM chip** with a **wide interface**
- **Flexibility:** memory controller does not need to deal with individual chips


- **Disadvantages:**

- **Granularity:** Accesses cannot be smaller than the interface width

DRAM Subsystem Organization


Dram: Dynamic RAM

DRAM Organization:




- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column

- Channel: Independent memory subsystem
 - E.g., 2 independent Channels:

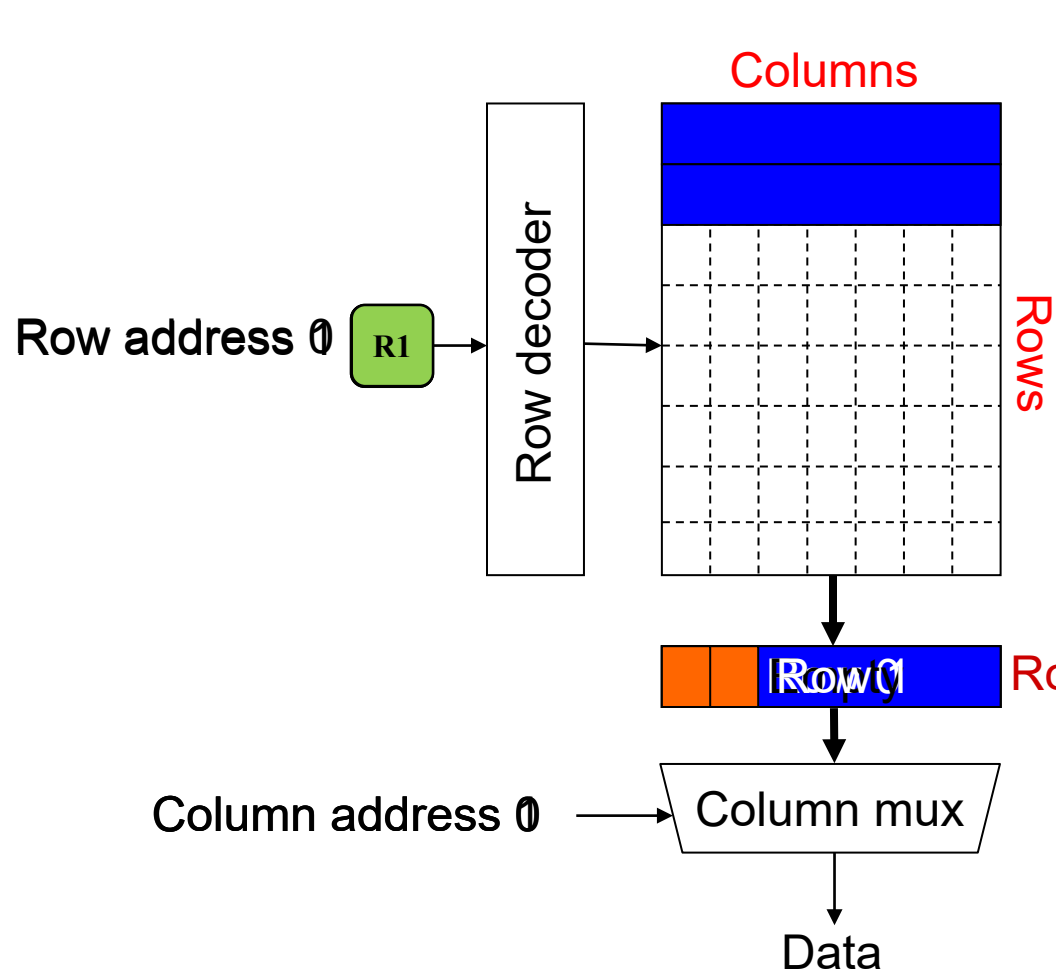


The DRAM subsystem



Each channel is connected to its own on-die memory controller

Row Operations & Row Buffer Locality



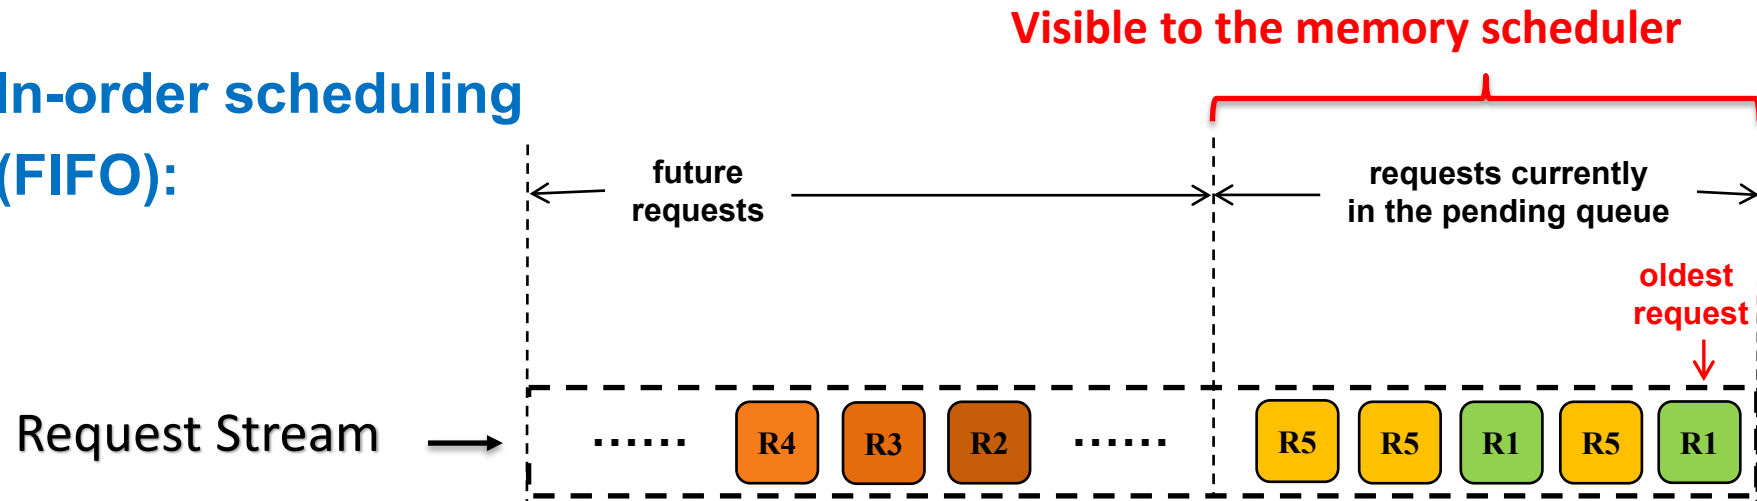
	Access Address:	Row Operation:
RBL=2	{ (Row 0, Column 0) (Row 0, Column 1)	Activation No operation
RBL=1	{ (Row 1, Column 0)	Restore, Precharge, Activation

CONFLICT !

Improving Row Buffer Locality (RBL) is the key to reduce DRAM energy consumption

RBL & Memory Scheduling Schemes


In-order scheduling (FIFO):



Activation Counter:

R1: Activation = 1
R5: Activation = 2
R1: Activation = 3
R5: Activation = 4
R5: Activation = 4 } Same activation
Avg RBL = 5 / 4 = 1.25

Out-of-order scheduling (FR-FCFS):



Activation Counter:

R1: Activation = 1 } Same activation
R1: Activation = 1 } Same activation
R5: Activation = 2 } Same activation
R5: Activation = 2 } Same activation
R5: Activation = 2 } Same activation
Avg RBL = 5 / 2 = 2.5

DRAM Milestones

	DRAM	Page DRAM	Page DRAM	Page DRAM	SDRAM	DDR SDRAM
Module Width	16b	16b	32b	64b	64b	64b
Year	1980	1983	1986	1993	1997	2000
Mb/chip	0.06	0.25	1	16	64	256
Die size (mm ²)	35	45	70	130	170	204
Pins/chip	16	16	18	20	54	66
BWidth (MB/s)	13	40	160	267	640	1600
Latency (nsec)	225	170	125	75	62	52

- In the time that the memory to processor **bandwidth** has more than **doubled** the memory **latency** has improved by a factor of only **1.2** to **1.4**
- To deliver such high bandwidth, the internal DRAM has to be organized as interleaved memory banks

Memory (DRAM) + Caches

- **It is important to match the cache characteristics**
 - Caches want information provided to them one **block** at a time (and a block is usually more than one word)
- **Memory design considerations:**
 - With the main memory characteristics
 - Use DRAMs that support fast multiple word accesses, preferably ones that **match the block size** of the cache
 - With the memory-bus characteristics
 - Make sure the memory-bus can support the DRAM access rates and patterns
 - With the goal of increasing the Memory-Bus-to-Cache bandwidth

Review: DRAM vs. SRAM

- **DRAM**


- Slower access (capacitor)
- Higher density (1T 1C cell)
- Lower cost
- Requires refresh (power, performance, circuitry)
- Manufacturing requires putting capacitor and logic together

- **SRAM**

- Faster access (no capacitor)
- Lower density (6T cell)
- Higher cost
- No need for refresh
- Manufacturing compatible with logic process (no capacitor)

Magnetic Disk

- **Purpose:** Long term, **nonvolatile** storage
 - Lowest level in the memory hierarchy: large, cheap, slow
- **General structure**
 - 1 to 4 rotating platter coated with a magnetic surface (2 sides recordable)
 - Rotational speeds of 5,400 to 15,000 RPM
 - Moveable read/write head to access the information for each platter
 - 10,000 to 50,000 **tracks** per surface
 - **Cylinder** - all the tracks under the head at a given point on all surfaces
 - 100 to 500 **sectors** per track
 - The smallest unit that can be read/written (typically 512B)
 - Outer tracks hold more sectors than the inner tracks



Magnetic Disk Characteristic

1. **Seek time:** position the head over the proper track

- 3 to 12/15 ms on average
- Due to locality of disk references, the actual average seek time may be only 25% to 33% of the advertised number


2. **Rotational latency:** wait for the desired sector to rotate under the head

- $\frac{1}{2}$ of $1/\text{RPM}$ converted to ms: $0.5R/5400\text{RPM} = 5.6\text{ms}$ to $0.5R/15000\text{RPM} = 2.0\text{ms}$

3. **Transfer time:** transfer a block of bits (one or more sectors) under the head to the disk controller's cache (70 to 125 MB/s are typical disk transfer rates)

- the disk controller's "cache" takes advantage of spatial locality in disk accesses
- cache transfer rates are much faster (e.g., 375 MB/s)

4. **Controller time:** the overhead the disk controller imposes in performing a disk I/O access (typically < 0.2 ms)



Typical Disk Access Time

The average time to read or write a 512B sector for a disk rotating at 15,000 RPM with average seek time of 4 ms, a 100MB/sec transfer rate, and a 0.2 ms controller overhead

Avg disk read/write


$$= 4.0 \text{ ms} + 0.5 / (15,000 \text{ RPM} / (60 \text{ sec/min})) + 0.5 \text{ KB} / (100 \text{ MB/sec}) + 0.2 \text{ ms}$$

$$= 4.0 + 2.0 + 0.005 + 0.2 = 6.2 \text{ ms}$$

If the measured average seek time is 25% of the advertised average seek time, then

$$\text{Avg disk read/write} = 1.0 + 2.0 + 0.005 + 0.2 = 3.2 \text{ ms}$$

Disk Latency & Bandwidth Improvement



- Disk **latency** = average seek time + rotational latency.
- Disk **bandwidth** is the peak transfer speed of formatted data from the media (not from the cache).

Flash Storage

- Flash memory is the first credible challenger to disks. It is semiconductor memory that is **nonvolatile** like disks but has latency 100 to 1000 times lower and is smaller, more power efficient, and more shock resistant.
 - Flash memory bits wear out. But with **wear leveling** it is unlikely that the write limits of the flash will be exceeded.
 - Example:

Storage	Write throughput	Read throughput	Latency
SSD	2,500 MB/S	3,500 MB/S	0.025 ms
DISK	250 MB/S	250 MB/S	5 ms

Flash Storage Characteristics


- **Flash memory is organized into blocks of pages**
 - page size 512B to 4KB, block size 32 to 128 pages (typical)
 - “read” is page read
 - “write” is block erase (~ 1 ms) followed by page write (and additional copying of other pages in that block)
- **Typically a type of Electrically Erasable PROGRAMMABLE Read-Only Memory (EEPROM)**
- **Comes in two flavors: NOR Flash and NAND Flash**
 - NOR Flash is randomly addressable (Minimum access size 512 bytes)
 - NAND Flash is less expensive (greater storage density) and is not randomly addressable (minimum access size 2048 bytes); so more popular

Dependability, Reliability, Availability

- **Reliability** – measured by the **mean time to failure (MTTF)**. **Service interruption** is measured by **mean time to repair (MTTR)**
- **Availability** – a measure of service accomplishment
$$\text{Availability} = \text{MTTF} / (\text{MTTF} + \text{MTTR})$$
- **To increase MTTF, either improve the quality of components or design the system to continue operating in the presence of faulty components**
 1. Fault avoidance: preventing fault occurrence by construction
 2. Fault tolerance: using redundancy to correct or bypass faulty components (hardware)

RAID: Redundant Array of Independent Disks


- **Arrays of independent physical disks working together as one logical disk**
 - Increase potential **throughput** by having many disk drives
 - Data can be spread across multiple disk
 - Multiple disk accesses can be made simultaneously for higher throughput
- **Reliability** for the array is lower than for a single disk
- **Availability** can be improved by adding redundant disks
 - Lost information can be reconstructed from redundant information
 - MTTR: mean time to repair is in the order of hours
 - MTTF: mean time to failure of disks is tens of years



RAID Disk Array


RAID: Level 0 (Striping, No Redundancy)

Assumes one stripe = four blocks



- **Multiple smaller disks as opposed to one big disk**
 - Multiple blocks can be accessed in parallel to increase the performance
 - Works well for large data requests
- **E.g., a 4-disk system gives four times the throughput (R/W) of a 1-disk system**
 - Same cost as one big disk – assuming 4 small disks cost the same as one big disk
- **What if one disk fails?**
 - No redundancy, data is lost
 - More likely to fail as the number of disks increases

RAID: Level 1 (Redundancy via Mirroring)



- **# redundant disks = # of data disks, so always two copies of the data**
 - Twice the cost of one big disk
 - Writes are made to both sets of disks and have no speed improvement
 - Reads can be 2 times faster
- **If a disk fails, the system just goes to the “**mirror**” for the data**

RAID: Design Choices


Category	Level	Disks	Features
Striping	0	N	No fault tolerance
Mirroring, no striping	1	2N	Expensive
Striping	2	N + m	Hamming code
Rarely used	3	N + 1	Parity disk
Block level Striping	4	N + 1	Parity disk
Parity	5	N + 1	Distributed parity
	6	N + 2	Dual distributed parity

- **Parity: used for error detection**
 - E.g., data: 00001111, parity bit: 0
- **Hamming code: limited error correction**
 - Using multiple parity bits to locate 1 error bit

How is the Hierarchy Managed?

- **registers ↔ memory**
 - by compiler (or programmer)
- **registers ↔ cache ↔ main memory**
 - by the cache & memory controller hardware
- **main memory ↔ external storage (flash, disk)**
 - Static: by the programmer with OS support (files)
 - Dynamic: by the operating system (virtual memory)
 - virtual address to physical address mapping
 - assisted by the hardware (TLB, page tables)


Virtual Memory



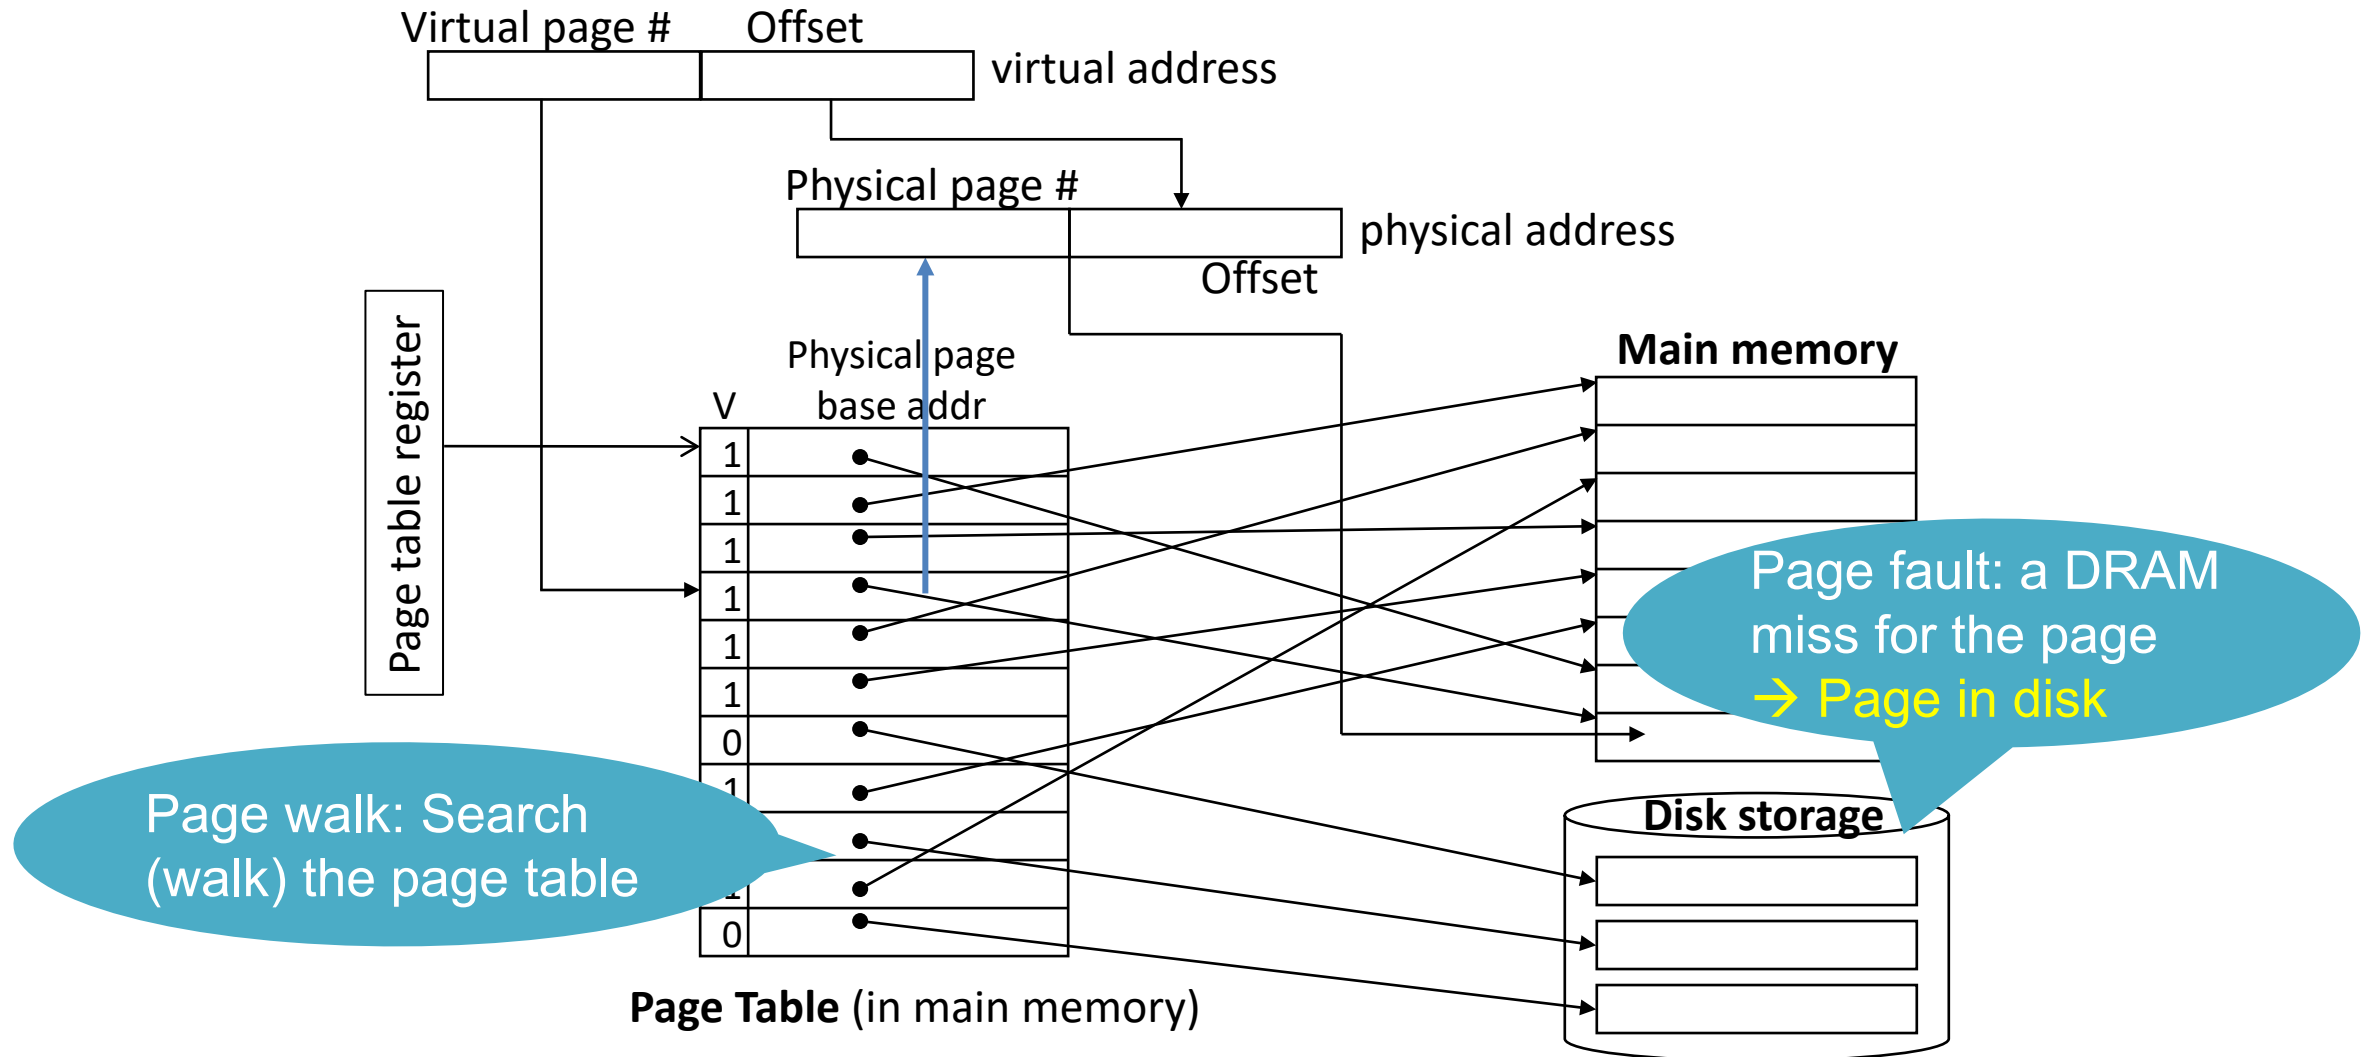
- **A virtual address must first be translated to a physical address to access memory**
 - Basic unit: page (e.g., 1KB to 64KB)
- **Virtual memory size can be larger than physical memory size**
 - Pages can be stored in the secondary storage

Address Space Isolation

- **Allows efficient and safe sharing of main memory among multiple processes**
 - The starting location of each page is contained in the program's page table
 - Create the illusion that each program has a large consecutive memory space
 - Improving memory utilization: code can be loaded anywhere the OS can find space




Address Translation Mechanisms




Virtual Addressing with a Cache

- It takes an extra memory access to translate a VA to a PA

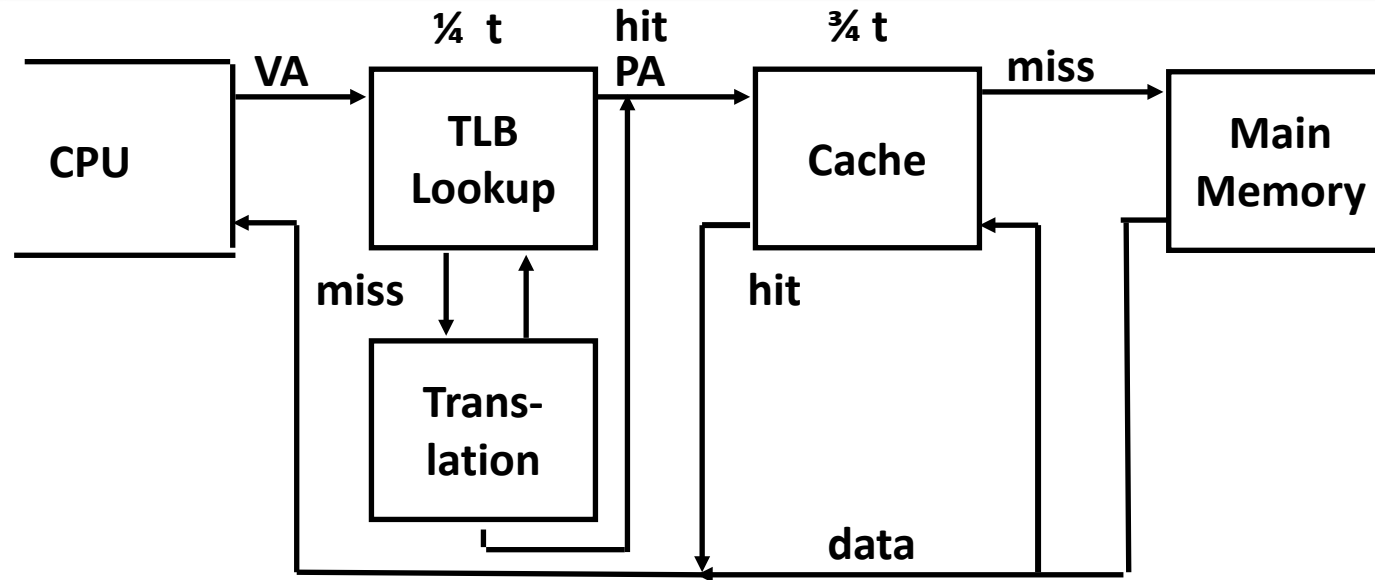


- This makes memory (cache) accesses very expensive
- The hardware solution is to use a Translation Lookaside Buffer (TLB)
 - A small cache that keeps track of recently used address mappings to avoid having to do a page table lookup in main memory

Making Address Translation Fast



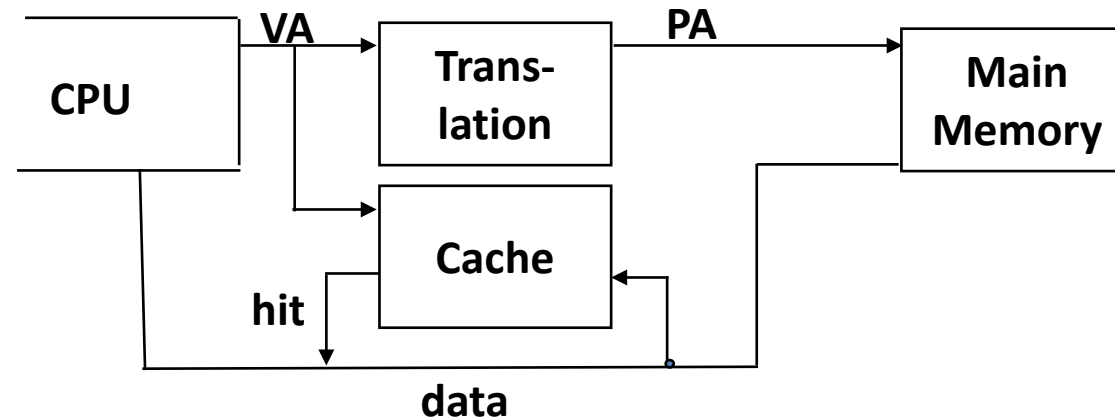
A TLB in the Memory Hierarchy



- **A TLB miss – is it a page fault or merely a TLB miss?**
 - If the page is loaded into main memory, then the TLB miss can be handled by loading the translation information from the page table into the TLB (10's of cycles)
 - If the page is not in main memory, then it's a true page fault (1,000,000's)
- **TLB misses are much more frequent than true page faults**

Why Not a Virtually Addressed Cache?

- A virtually addressed cache would only require address translation on cache misses

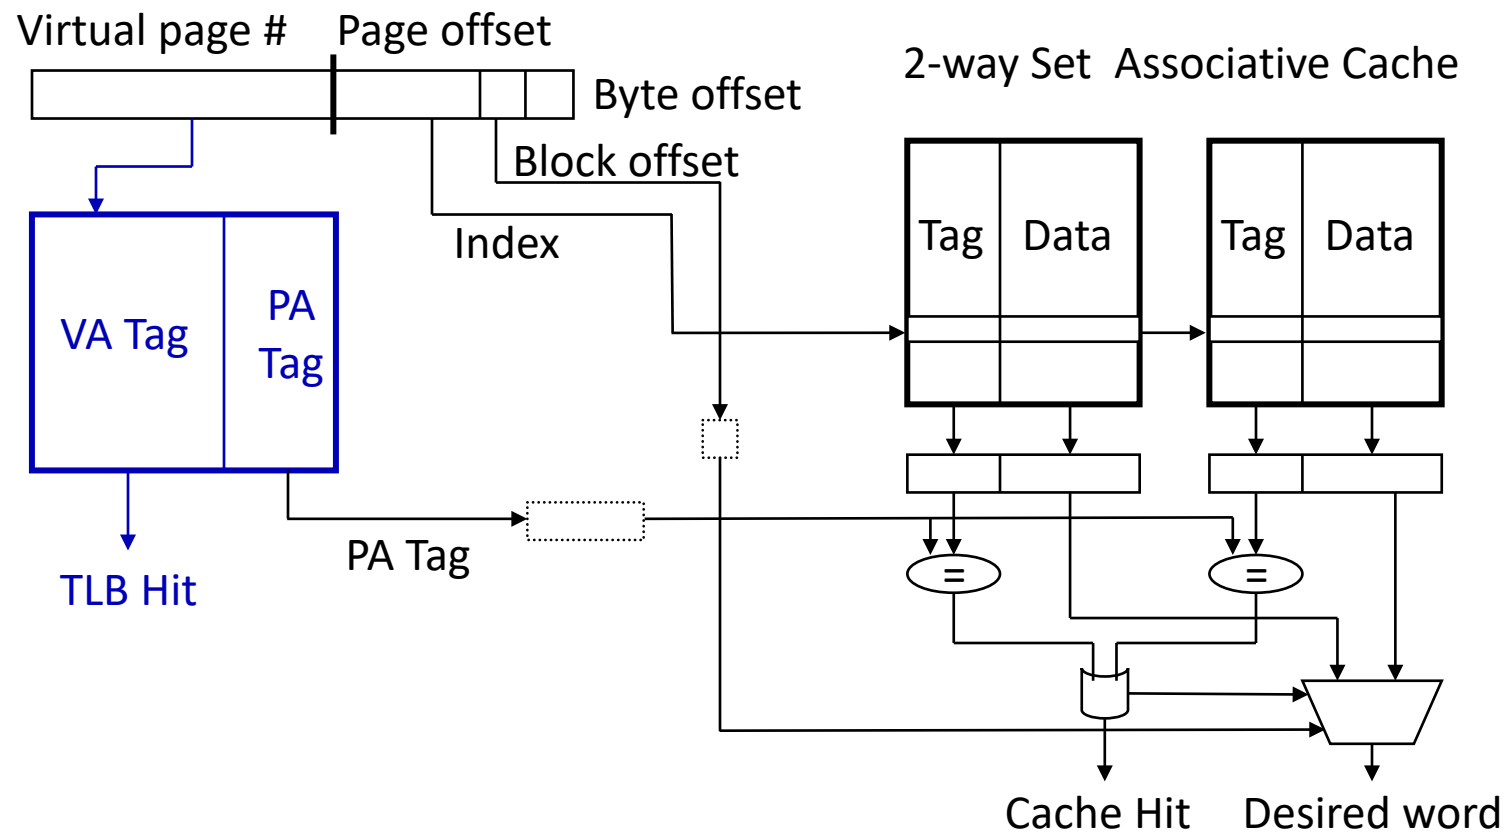


Aliasing issue:

- Two programs which are sharing data will have two different VA for the same PA
 - Two copies of the data in cache and two entries in the TLB → coherence issues
 - Must update all cache entries with the same physical address or the memory becomes inconsistent
- No address space isolation: data can be corrupted

Reducing Translation Time

- **Overlapping the cache access with the TLB access**
 - Higher bits for TLB; lower bits for cache.



SAN JOSÉ STATE UNIVERSITY *powering* SILICON VALLEY

