

CMPE 200  
Computer Architecture & Design

## **Lecture 5.** **Parallel Processing (2)**




Haonan Wang



SAN JOSÉ STATE  
UNIVERSITY

# Multiprocessor


- **Parallelism in processor unit → Multiprocessor**
  - Multiple CPUs or Multiple Cores within a CPU
- **Two representative types:**
  - Shared-memory
  - Message-passing
- **Platform variety: #cores, hierarchy, #memory controllers, interconnect, ...**



# Multicore Power Management

- For power management, you can turn off whole cores, vary the core frequency (lower or higher), and run fewer threads.
- **Power reduction modes**
  - Nap: optimal wakeup latency relative to power reduction
    - Clocks turned off to the execution units, core frequency reduced, caches and TLB's still active, remain coherent
  - Sleep: optimal power reduction relative to wakeup latency
    - Caches and TLB's are purged, all clocks turned off, voltage lowered, some state info is retained

4 PowerPC Architected States



# Multiprocessor Key Questions

---

- **Q1 – How do they share data?**
- **Q2 – How do they coordinate?**
- **Q3 – How scalable is the architecture? How many processors can be supported?**

# Multiprocessor Key Questions

---

- **Q1 – How do they share data?**
- Q2 – How do they coordinate?
- Q3 – How scalable is the architecture? How many processors can be supported?

# How to Share Data?

- **Shared Memory multi-Processor (SMP)**


- Single address space shared by all cores
- Cores coordinate/communicate through shared variables (via loads and stores)
  - Need [synchronization](#) primitives (e.g., locks)
- Two styles
  - Uniform memory access ([UMA](#)): easier for programming
  - Non-uniform memory access ([NUMA](#)): more scalable & lower latency to local memory

- **Message Passing multi-Processor (MPP)**

- Each core has its own private address space
  - Cores share data by *explicitly* sending and receiving information ([message passing](#))
- Coordination is built into MP primitives ([message send](#) & [message receive](#))

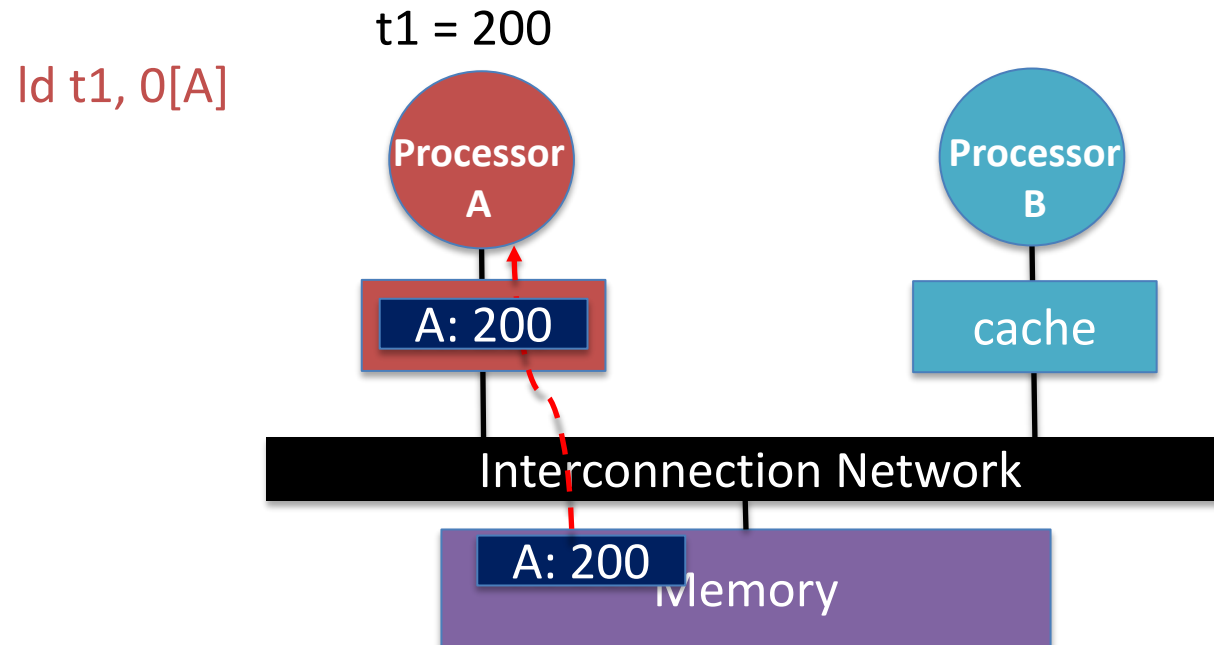
# Shared Memory Model

- **Parallel programs communicate through shared memory**
  - E.g., processor A writes to an address, then processor B reads from the same address
- **Each read should receive the most up-to-date value**



# Shared Data Management

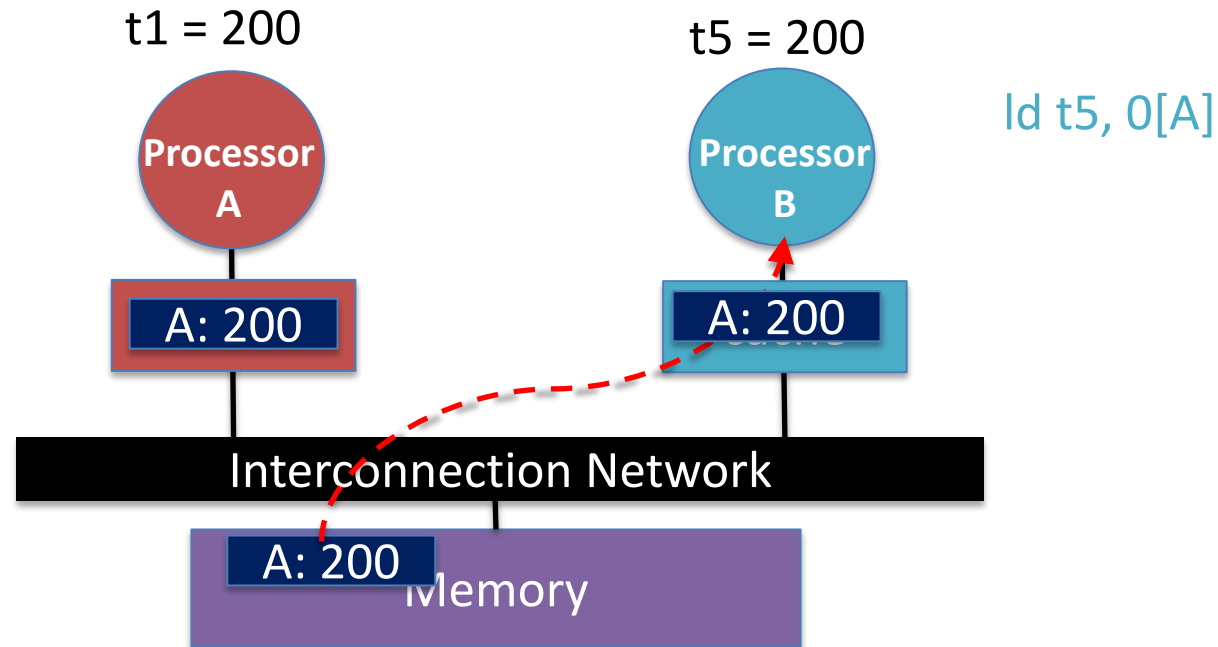
- If multiple processors cache the same block in their private cache, how do they ensure they all see a consistent state?





# Shared Data Management

- If multiple processors cache the same block in their private cache, how do they ensure they all see a consistent state?

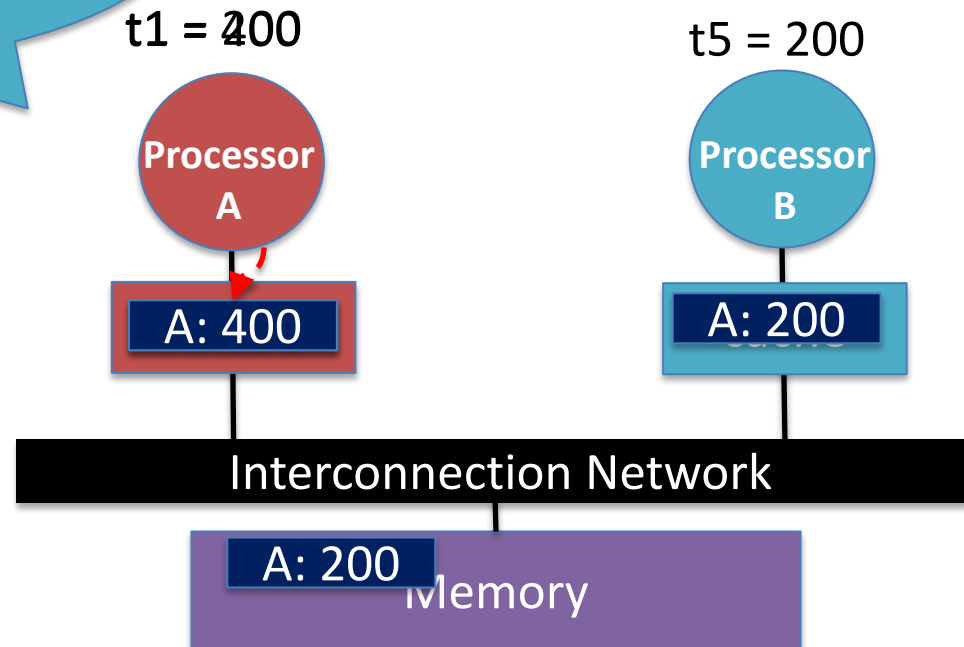


# Shared Data Management

- If multiple processors cache the same block in their private cache, how do they ensure they all see a consistent state?

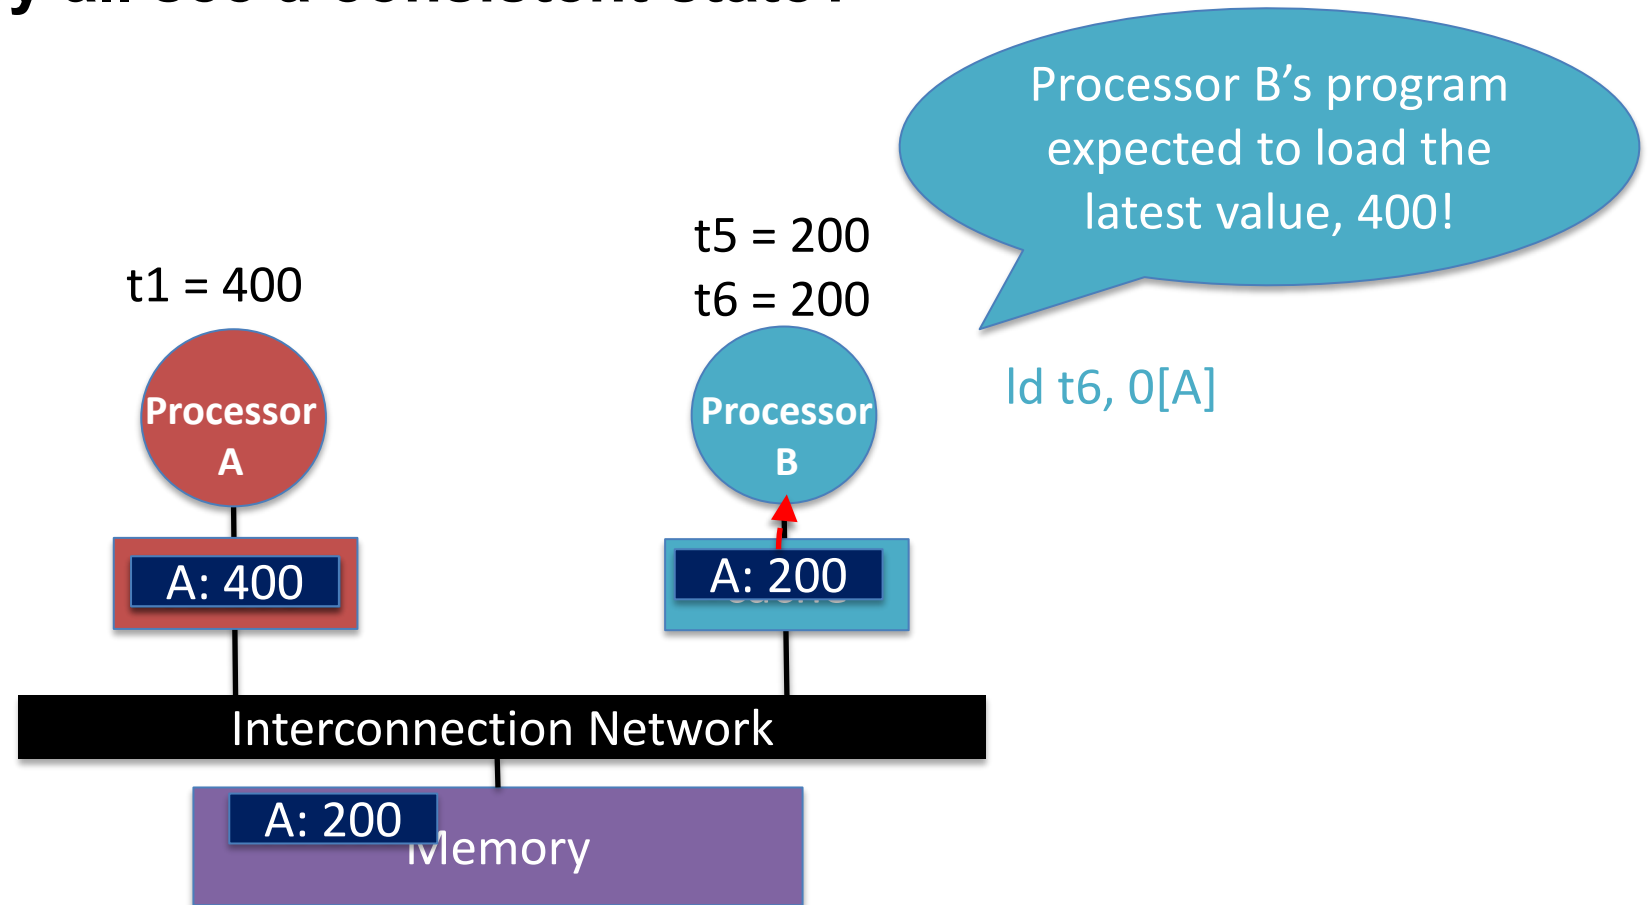
Writeback cache updates memory only when the block is replaced

addi t1, t1, #200  
st t1, 0[A]



# Shared Data Management

- If multiple processors cache the same block in their private cache, how do they ensure they all see a consistent state?




# A Coherent Memory System

- **Coherence:** writes to the same location are seen in the same order by all cores (serialized)
- **Cache coherent multiprocessor:** provide **migration & replication**
  - Migration: data items are moved to local caches and used in a transparent fashion
    - Reduces latency and the bandwidth demand on the shared memory
  - Replication: simultaneous reads make copies of the data item in local caches
    - Reduces both latency and contention for a **read** shared data item

# Cache Coherence Methods

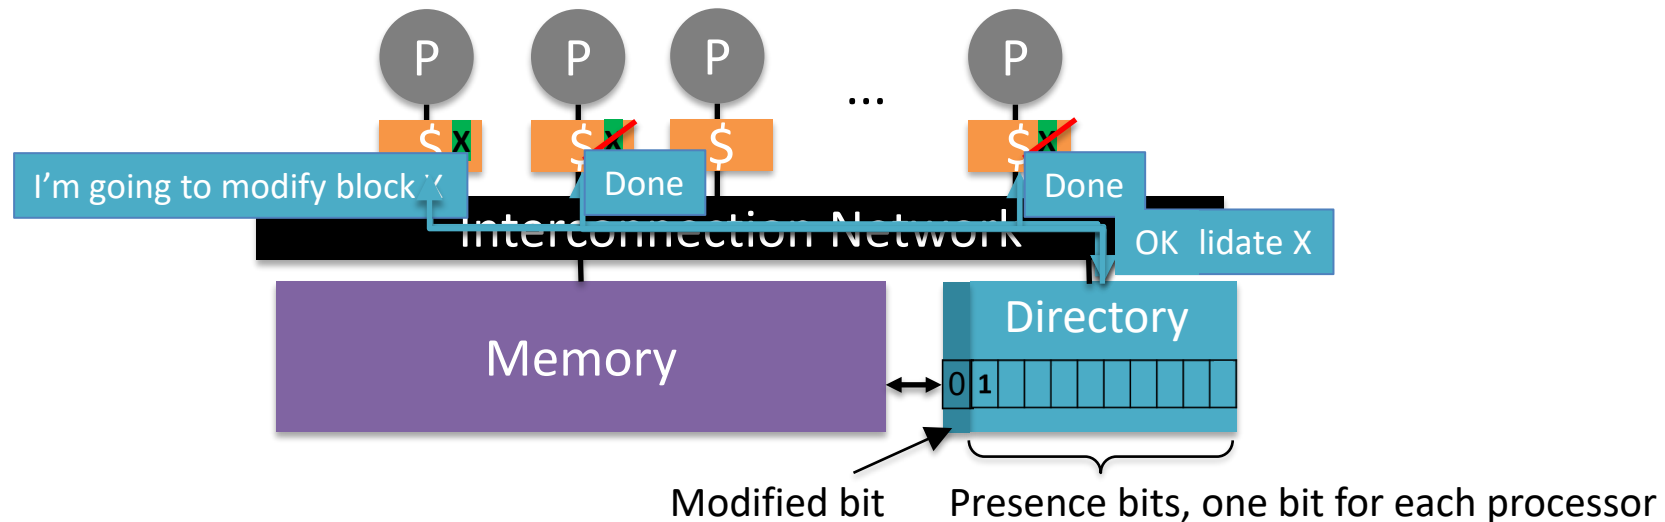
- **Snooping-based protocol**
  - Each processor's cache controller constantly snoops on the bus
  - Processors observe other processors' actions
    - E.g., processor A makes “read-exclusive” request for A on bus, processor B sees this and invalidates its own copy of A



# Cache Coherence Methods

- **Directory-based protocol**


- Directory tracks ownership (sharer set) for each block (who has what)
  - A *modified* bit and multiple *presence* bits per cache block
- Directory coordinates invalidation appropriately
  - E.g., processor A asks directory for “read-exclusive” copy, directory asks processor B to invalidate the requested copy, waits for ACK from processor B, then responds to processor A



# Write Invalidate (All Other Copies)

- **Can enforce write serialization**

- Two processors attempt to write the same data at the same time → one of them wins the race → the other core's copy is invalidated
- The other core must obtain a new copy of the data → must now contain the updated value




- **False sharing problem**

- Two cores writes to two different variables in the same cache block → increases miss rate
- Compilers can help reduce this by allocating correlated data to the same cache block

# Message Passing multi-Processors (MPP)

- **Each core has its own private address space**
  - Cores share data by explicitly sending and receiving information (**message passing**)
- **Coordination is built into message passing primitives (**message send** and **message receive**)**





# Pros and Cons of Message Passing

- **Cons:**

- Performance
- Harder to port a sequential program to an MPP
  - Every communication must be identified in advance

- **Pros:**

- Easier for hardware design
- Communication is explicit, fewer “performance surprises”
  - Message passing standard: MPI


# Multiprocessor Key Questions

---

- Q1 – How do they share data?
- **Q2 – How do they coordinate?**
- Q3 – How scalable is the architecture? How many processors can be supported?

# SMP Example with 10 Cores

sum[C0] sum[C1] sum[C2] sum[C3] sum[C4] sum[C5] sum[C6] sum[C7] sum[C8] sum[C9]




# Process Synchronization

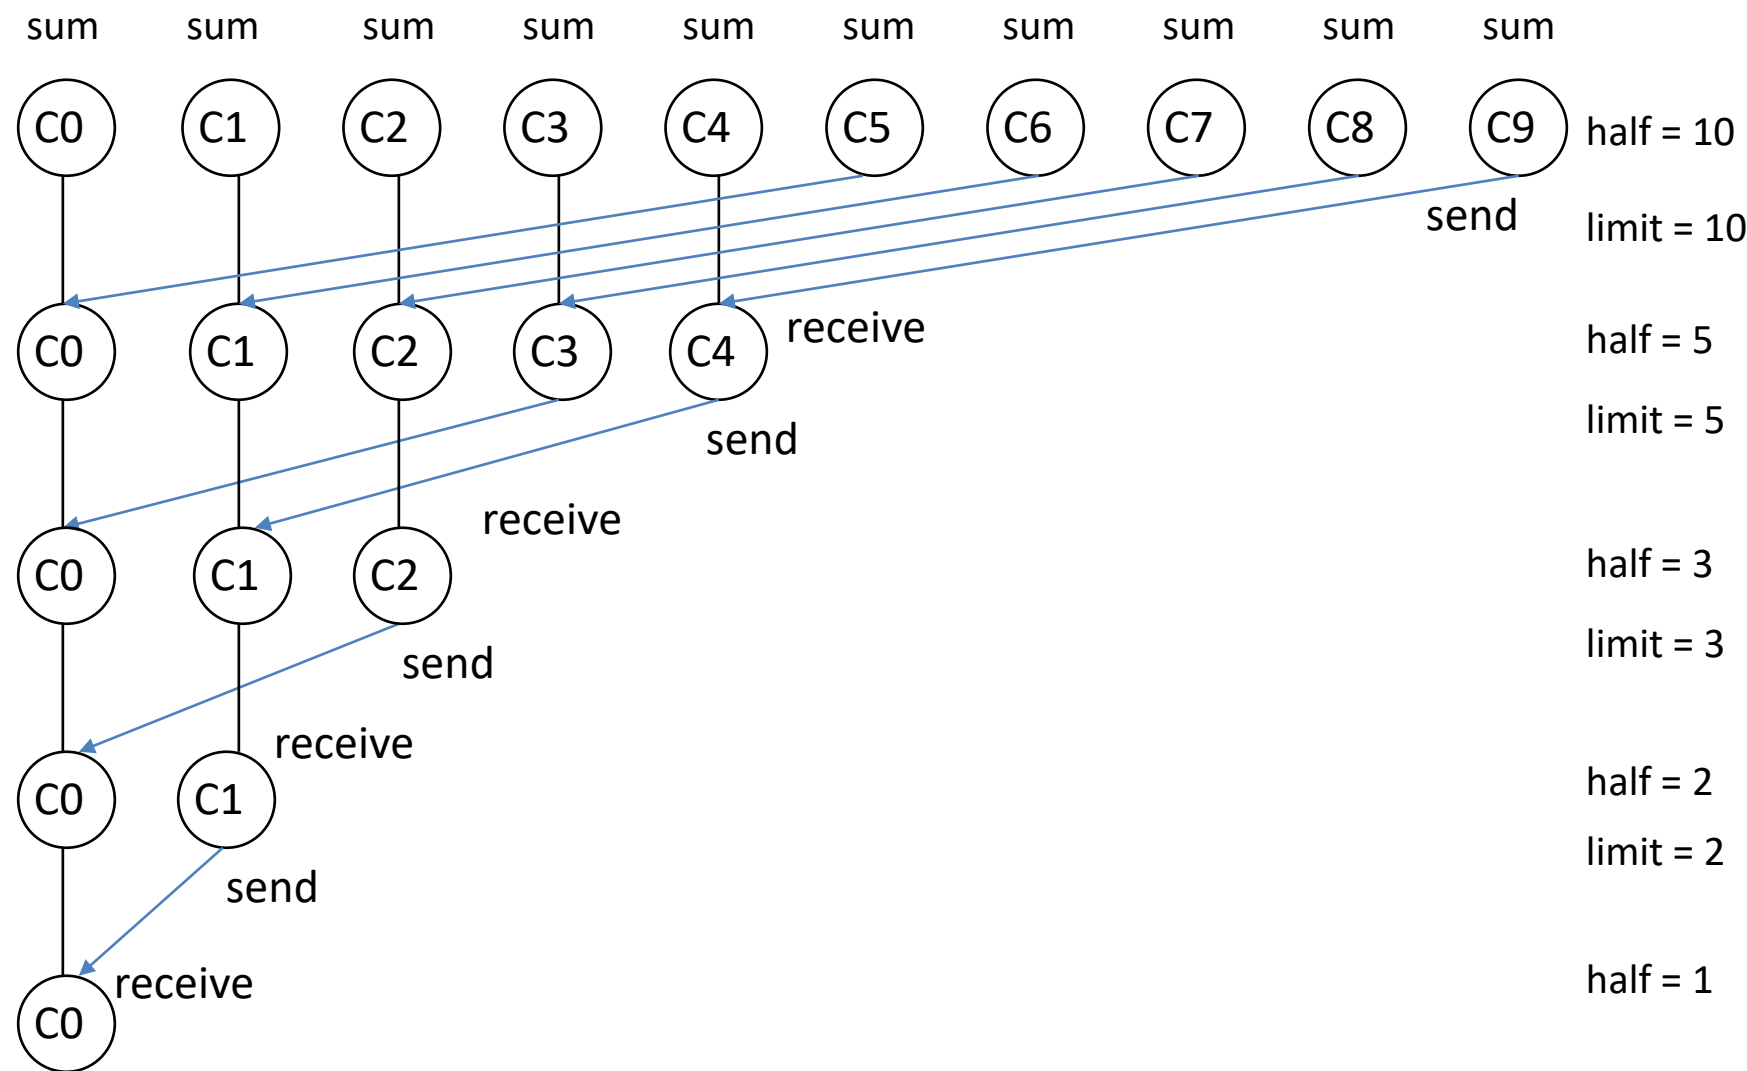
- **We need to coordinate processes working on a common task**
  - Locks (e.g., spin locks, semaphores)
    - **Mutual exclusion**: restrict data access to one core at a time
    - **Sequential ordering**: must complete the first operation before starting the second
- **Need architectural support to implement the lock**
  1. Arbitration mechanism
    - Single bus provides an arbitration mechanism, since the bus is the only path to memory  
→ the core that gets the bus wins
  2. Architecture-supported operation
    - **Atomic swap operation** (e.g., `ll` and `sc` on MIPS) allows a core to both read a location and set it to the locked state (**test-and-set**) in the same bus operation

# Process Synchronization

- **synch(): Cores must synchronize before the “consumer” core tries to read the results from the memory location written by the “producer” core**
  - Barrier synchronization: a synchronization scheme where cores wait at the barrier, not proceeding until every core has reached it



# MPP Example with 10 Cores



# Load Balancing

- **Load balancing is important for multiprocessor performance**
  - E.g., a single core with twice the load of the others cuts the speedup almost in half
- **S = serial part, W = parallel part, N cores**
  - Time on 1 processor is  $S + W$
- **Suppose the workload is equally distributed and there is no extra overhead to accumulate partial results**
  - Time on N processors is  $S + W/N$
- **Suppose the workload is almost equally distributed (and still no overhead)**
  - One core does  $2(W/N)$ , one does nothing, most do  $W/N$
  - Time is  $S + \max(2(W/N), 0, W/N) = S + 2(W/N) = S + W/(N/2)$

# Multiprocessor Key Questions

- Q1 – How do they share data?
- Q2 – How do they coordinate?
- **Q3 – How scalable is the architecture? How many processors can be supported?**

			# of Cores
Communication model	Message passing		8 to 2048
	SMP	NUMA	8 to 256
		UMA	2 to 64
Physical connection	Network		8 to 256
	Bus		2 to 36



# Conclusion Time

---

What is the difference between concurrency and parallelism?

Software implementation vs. hardware capability

What is the main difference between UMA and NUMA?

Memory access latency

SAN JOSÉ STATE UNIVERSITY *powering* SILICON VALLEY

